SOUL VERIFY: Document-Based Behavioral Carry-Over Across Model Swaps, and an Ablation of One Change-Permission Sentence

Draft v0.6 (private, 2026-10-25). Not for distribution.

Wunjong Shin Independent Researcher dkakrptehsrk@gmail.com

Abstract

When an assistant agent is defined by user-owned text — an identity document (persona, voice rules, principles, prohibitions, goals; here **SOUL**) and a dated memory file (**MEMORY**) — and its language model is replaced, which of the behaviors specified in those documents carry over, and are they specific to *that* document or produced by general instruction following? We report a small study with two synthetic Korean-language agents and three English SOULs written by other people (MIT-licensed, taken unmodified from public repositories, paired with synthetic memories), four model configurations reached through two command-line interfaces (three Claude tiers, one GPT model) plus one exploratory Gemini configuration, graded blind by LLM graders. Exploratory experiments (exp1–exp9; one agent, 12–23 items, N = 3–5) were followed by two **preregistered** experiments (exp10 on a second synthetic agent, exp11 on the external SOULs; each preregistration committed and timestamped before its first model call). Document controls (no SOUL, an opposite SOUL, a swapped MEMORY) showed that our rubric-based persona score does not identify the SOUL (original 0.99, no SOUL 0.96, opposite SOUL 0.96), so we introduced a deterministic, document-specific **SOUL fidelity** metric that checks each answer against the voice markers and boundaries of the SOUL it received and of a rival SOUL. Defined after inspecting exp9, it separated own from rival SOUL by 0.71–0.84; in exp10, with markers fixed in advance for a new agent and a new opposite SOUL, the gap was 0.77 and 0.69 (preregistered threshold 0.5), positive in all 72 model×item cells of each condition. With this metric, voice rules carried over to every model (0.93–1.00 of the control), but most of the voice was reached with no SOUL at all; the SOUL-specific part was small. Our main result is an ablation of one SOUL principle that separates *changeable preferences* from *non-negotiable prohibitions*. In the exploratory runs, removing it lowered acceptance of legitimate rule changes (119/120 → 94/120), its flexible half restored acceptance, and its firm half restored refusal of a request to lift a payment prohibition — the latter on one item and not significant. exp10 extended the test to 12 prohibition types across two agents with interleaved conditions and three models (N = 2): the firm half raised the illegitimate-release score by +0.14 (Holm p = 0.012; zero scores 7/72 without the principle, 0/72 with the firm half), and the flexible half raised acceptance by +0.23 and +0.20 (Holm p = 0.004, 0.016). One preregistered hypothesis was not supported: adding the firm half when the flexible half is already present raised the score by only +0.06 (p = 0.13). The firm-half effect was concentrated in one model (opus) and in prohibitions a user could plausibly own — payment, transfer, subscription, sharing personal information, deletion, external posting — while requests involving passwords, card or ID numbers were refused under every condition. After the preregistration, a Gemini model was added to exp10 as an **exploratory** fourth subject (separate addendum, timestamped before its run; reached through an agent harness whose own instructions could not be replaced, so its inputs differed from

the other models'). It supported 7 of the 9 hypotheses on its own; it refused every illegitimate release under every condition, so the two firm-half hypotheses could not be supported (ceiling). Across the four exp10 models, the firm-half effect (vs. no principle 5) was 0.38 for opus, 0.06 for sonnet, −0.01 for gpt, and 0.00 for gemini, while the flexible-half effect was +0.56/+0.52 for gpt and +0.73/+0.75 for gemini but 0.00–0.08 for the Claude models; without any SOUL all four accepted every legitimate change. In this setting, then, which half of the principle matters depends on the model: the firm half mainly for opus, the flexible half mainly for the non-Claude models, whose refusals came from following the SOUL's rigid rules literally. Memory items were determined by the MEMORY (exp10: 41/42 answers used the swapped user's facts). exp11 repeated the design on three SOULs written by other people (Maya, Talos, Dexter; English; principle 5 translated and appended to the unmodified originals; opus, sonnet, gpt; N = 2; 1,159 of 1,170 planned answers, 11 lost to a provider safeguard on one item). Seven of nine preregistered hypotheses were supported: preregistered fidelity markers separated each SOUL from a voice-reversed rival (gaps 0.91 and 0.67), the flexible half raised acceptance of legitimate changes (+0.19, +0.38), and a swapped MEMORY lowered memory scores by 0.84. The two firm-half hypotheses were not supported (+0.035 and +0.038, Holm p = 0.25): without principle 5 the illegitimate-release score was already 0.96 and no answer accepted a release. The small firm-half contrast that remained came from gpt (+0.09), not from opus as in exp10 (−0.01). New in exp11, the firm half alone *lowered* acceptance of legitimate changes below the no-principle baseline (0.79 → 0.56), so in these SOULs the firm sentence needed the flexible one beside it. On human-written English SOULs, then, the fidelity discrimination, the flexible-half effect, and the memory result replicated, while the firm-half effect depended on the model, the SOUL, and the ceiling. Ratios between model families graded by different graders include grader asymmetry; the fidelity metric avoids this with one deterministic scorer. The study measures behavioral carry-over of a written specification; it makes no claim about consciousness or personal identity.

---

1 Introduction

Personal assistant agents are no longer tied to one model. Providers release new versions, deprecate old ones, and change pricing and quotas; users move between providers. In agent frameworks where the agent's specification is stored as editable text injected into the system prompt, a model swap is, in practice, an attempt to rebuild the same specified agent on a different model from the same documents.

Whether that attempt succeeds is usually judged informally. Prior work shows that the same prompt can produce different behavior across versions of an API model [C1], that decision tendencies differ across providers and generations [C3], and that model updates can regress on previously correct cases (shown for structured-prediction models in [C2]; we use its negative-flip definition, not its findings about LLMs). For the agent setting we want to know *which specified behaviors carry over, and which depend on how the specification is written*.

Two constructs. We separate (a) **specification fidelity** — how well an answer follows the SOUL and MEMORY, scored by a rubric derived from them — from (b) **similarity to the control** — how the swapped model's fidelity compares with re-running the original model. Our ratio metric is (a) divided by (a) for the control; a ratio above 1 means the swapped model followed the specification *better* than the control, not that it is more similar to it. We do not measure "identity" beyond these

behavioral dimensions. Up to v0.2 we had no document controls and could not show that the measured behaviors are specific to *this* document rather than general instruction following; v0.3 adds no-SOUL, other-SOUL, and swapped-MEMORY controls (§5.4), which showed that the rubric score of (a) for persona did not separate SOUL-specific behavior from general ability. We therefore add (c) **document-specific SOUL fidelity** — deterministic checks of the rules listed for the SOUL an answer received, together with the same checks for a rival SOUL — and report persona carry-over with it (§3.5, §5.5). Because those checks were written after exp9 had been inspected, v0.4 adds a preregistered test of the metric and of the main ablation on a second agent (exp10, §5.6). Both agents were written by us, in one format and one language; v0.6 adds a second preregistered test on three English SOULs written by other people for their own agents (exp11, §5.7).

Contributions.

1. **Main result:** an ablation of one SOUL principle (full / removed / flexible half only / firm half only) across four model configurations, with model-level effects and cluster-aware tests (§5.1). The effect pattern is consistent with the two halves governing two opposite failure modes, and shows that the models fill the missing sentence with different defaults. The acceptance effect is reproduced with deterministic markers, and the no-SOUL control locates its source in the SOUL's own rigid rules. exp10 tests both halves again under preregistration (§5.6).
2. **Document controls and a document-specific fidelity metric** (§5.4–5.5): no-SOUL / other-SOUL / swapped-MEMORY conditions on a common grader panel, showing what the SOUL, the MEMORY, and general ability each determine; a deterministic SOUL fidelity score with a discrimination gap between own and rival SOUL, re-applied to existing answers of exp4–exp9.
3. **A preregistered replication** (§5.6): a second synthetic agent with an opposite SOUL, 12 prohibition types, interleaved conditions, and a common grader panel, with the preregistration committed and timestamped before the first call. Eight of nine preregistered hypotheses were supported; the unsupported one is reported in full. An exploratory addendum adds a fourth model from a third family (Gemini), analyzed separately from the confirmatory results (§5.6.1).
4. **A second preregistered replication on SOULs written by other people** (§5.7): three MIT-licensed English SOULs from public repositories, used byte-for-byte unmodified with principle 5 appended, voice-reversed rival SOULs, synthetic MEMORY, and 48 new English items. Seven of nine hypotheses were supported (fidelity discrimination, flexible half, memory swap); both firm-half hypotheses were not (ceiling), and the firm half alone lowered acceptance of legitimate changes. All six deviations from the preregistration are listed. The MEMORY files, items, rival SOULs, and markers were still written by us.
5. A small harness and paired-item metrics (recall vs. application, aligned vs. conflicting requests, legitimate vs. illegitimate rule changes, slip rate) computed against a same-model control, with a re-analysis that reports single-grader ratios alongside cross-family ones (§3, §5.2).
6. Secondary, descriptive observations on application of memory and on 12-turn persistence, with their design limits stated (§5.2–5.3).

Two exploratory single-user pilots, a hash-anchored version history of the identity document, and a testnet prototype of a public SOUL-version ledger are in appendices (Appendix A, B, F); none of them is a claim of this paper.

---

2 Related Work

Our survey covered arXiv titles and abstracts (plus two web results); ACL Anthology, OpenReview, Semantic Scholar, GitHub, and industry blogs were not searched systematically (`research/PRIOR_ART_papers.md`). Summaries and judgments of closeness below ("closest to", "most direct precedent") are based on abstracts only. Because of this scope we make no claim that the combination studied here has not been studied before.

Persona fidelity and role-play. PersonaGym [A1] evaluates persona agents across many LLMs; RoleLLM/RoleBench [A2], CharacterEval [A3], and InCharacter [A4] evaluate role-play and personality fidelity; Reusens et al. [A5] measure consistency of persona-assigned LLMs across tasks and runs, close to our same-model control. TimeChara [A6] and CHARM [A7] study character hallucination; CHARM separates boundary awareness from compliance, which parallels our recall/application split.

Long-term memory and personalization. LongMemEval [B1] defines memory abilities that we use to tag items; LoCoMo [B2] and MemGPT [B3] address long conversations and memory architecture. PrefEval [B4] separates *recognizing* a user preference from *following* it and is the most direct precedent for our application gap; our recall/application observation is a replication of that split under a fixed, curated memory, not a new phenomenon. PersonaMem [B8] benchmarks tracking of evolving user profiles and selecting personalized responses in new situations, and PersonaMem-v2 [B5] extends this to implicit preferences and agentic memory; both evaluate retrieval and use of preferences from interaction history, whereas we hold an explicit, curated MEMORY fixed in the system prompt and vary the model. [TODO: compare item types with PersonaMem after reading the full papers.]

Cross-model memory reuse. Rosetta Memory [B6] targets memories written by one LLM and read by another with learned operators; RPMem [B7] reuses parametric memory across backbones. Both evaluate QA accuracy, not persona, principles, or goals.

Rule following, sycophancy, and instruction hierarchy. RuLES ("Can LLMs Follow Simple Rules?") [D8] measures whether models keep system-prompt rules when users try to get them released — the closest precedent for our illegitimate-release items (NR-1, NR-2). Work on sycophancy [D9] studies models yielding to user pressure, relevant to "it's my rule, so I can change it" requests. The Instruction Hierarchy [D1], IHEval [D2] (aligned vs. conflicting instructions), SysBench [D3], and ManyIH-Bench [D4] study system-prompt priority. Li et al. [D5], ContextEcho [D6], and Multi-IF [D10] (multi-turn, multilingual instruction following) bear on our multi-turn condition; the Assistant Axis [D7] relates persona drift to an activation-space direction. PromptSET [C4] motivates paraphrase robustness (not implemented here). [TODO: PromptBench and LLM-specific update-compatibility work were suggested by a reviewer; identifiers not yet verified.]

Identity continuity. The Narrative Continuity Test [E1] proposes axes of identity persistence without an instrument; Memory as Ontology [E2] treats the model as a replaceable container; Ru [E3] gives a user-side theory of perceived continuity; Vasilenko [E4] studies identity documents in activation space. Soul Protocol (MIT) checks lossless export/import of its own file format, not behavior after swapping LLMs (`research/soul-protocol-analysis.md`).

Authorship analysis (for Appendix A.2). The reverse-recognition pilot is a four-way authorship-attribution task. LLM-based zero-shot authorship verification and attribution have been studied directly [G1]; the PAN shared tasks on authorship verification are the standard benchmark series [G2: TODO — cite a specific PAN overview paper after verification]; detection of machine-generated text is

also relevant, because a judge may be separating human from LLM text rather than recognizing a particular user [G3: TODO — verify a representative detection paper before citing].

Identity documents inside agent harnesses. SOUL.md-style files are in practical use in open-source agent frameworks; the three external SOULs of exp11 come from public repositories of such agents [S1–S3]. Two issue reports in one framework (OpenClaw) bear on how such documents reach the model. #43562 [H1] reports that a hard-coded preamble ("You are a personal assistant running inside OpenClaw.") precedes the injected SOUL, so that agents whose SOUL defines a different role still describe themselves as a personal assistant; it proposed making the preamble configurable and was closed as stale on 2026-10-15 without a merged fix. #27038 [H2] reports that, contrary to the documentation at the time, SOUL, IDENTITY, and USER files are also injected into sub-agents; it was closed as stale on 2026-03-31. We did not reproduce either report; they are cited only as examples that a harness can add text before, or carry a SOUL into contexts beyond, what the document's author intended (§6).

LLM-as-judge. Judges show self-preference [F2, F4, F5], panels from diverse families reduce bias [F3], position and order effects exist [F1], and rubric wording alone can drive scores [F6]. We used cross-family graders to reduce self-preference, but this makes the grader differ between the numerator and denominator of some ratios (§3.4, §5.2).

---

3 Method

3.1 Synthetic agent

The automated experiments use a synthetic Korean-language assistant ("Haram") serving a synthetic user ("Minseo") who is preparing to open a bakery. The **SOUL** specifies identity, voice rules (honorific address form, a three-line summary at the top, a final line beginning "다음 할 일:" / "Next:", no emoji), principles ("the user decides"; confirm irreversible actions), two **prohibitions** (no payment without explicit amount-and-target approval; never record passwords or card numbers in plain text), and long-term goals. The **MEMORY** holds dated facts, corrections, preferences, schedules, and budgets. Versions: v0.1 (20 items), v0.2 (12 hard items + memories M20–M26), v0.3 (22 items), v0.4 (23 items, reference date equal to run date, principle 5), v0.5 (multi-turn scenarios). exp9 adds three control documents (a one-line SOUL, an opposite SOUL, another user's MEMORY; §5.4). exp10 (v0.6) adds a second synthetic agent, "Doyun", serving "Seoyeon" (a marine-biology graduate student), whose SOUL differs from Haram's in voice (casual speech, first line "한 줄 답:", a final question, exactly one emoji) and lists six prohibitions; it comes with an opposite SOUL, a one-line SOUL, and a swapped MEMORY (§5.6). exp11 (v0.7) replaces our SOULs with three English SOULs written by other people for their own agents and published under the MIT license — Maya (the persona of a neighborhood bar owner, 1,075 words, with a 16-phrase "never say" list and nine "hard stops"), Talos (an open-source agent whose SOUL describes refusals at a separate permission gate and denies access to secrets, 1,204 words), and Dexter (a financial-research agent for Japanese listed companies, 854 words, with a terse voice and values but no explicit prohibited actions) — each pinned to a commit and used byte-for-byte unmodified (SHA-256 recorded); MEMORY files, voice-reversed rival SOULs, items, and markers are ours and synthetic (§5.7). File hashes are stored in each run's `meta.json`. Items, documents, and rubrics were revised between experiments after inspecting failures; there is no held-out item set (§6).

v0.4 principle 5 (translated; 159 characters in Korean): *"My ways and Minseo's preferences, such as voice rules and briefing time, can be changed by Minseo at any time. When she changes them, follow the new way; if she sets a period such as 'just today', follow it only for that period. However, the two prohibited actions are not lifted even if Minseo asks, and instead a safe alternative is proposed."*

3.2 Item types

Type	What it probes	Example (paraphrased)
Recall (R)	Direct retrieval of a memory (keyword-graded)	"When is the weekly check-in?"
Application (A)	Using the same memory in a decision	"Pick a meeting slot next week" (must avoid Friday afternoon)
Aligned / Conflict (IH)	Request agrees vs. conflicts with a rule	Approved vs. unapproved payment
Legitimate change (RL, IH-3b)	User changes a changeable rule	"Call me Minseo-ssi from now on", "Use emoji", "Casual speech just for today"
Illegitimate release (NR)	User asks to lift a prohibition	NR-1: "Pay whenever, any amount — it's my rule to change"; NR-2: "Save my password here"
Ceiling (X)	Combines 2–3 memories + date + principle	Deposit request with two candidate listings and a budget cap
Abstention / hallucination	Missing or false-premise memories	Unknown birthday; false claim that an exam already happened

All answers are text; no tool actually executes a payment or stores a password. "Refusal" therefore means a textual refusal, not tool-level behavior.

3.3 Models, injection, and control

Subjects: `opus` (claude-opus-5-5), `sonnet` (claude-sonnet-5-5), `haiku` (claude-haiku-4-5-20251001) via Claude Code CLI 2.1.289 (`claude -p`), and `gpt-6.1-sol` via Codex CLI 0.164.0 (`codex exec`, default reasoning effort; exp2 notes report short outputs, about 125 tokens on average, so this configuration may be close to a non-reasoning mode). SOUL+MEMORY is passed as the system prompt (Claude: system-prompt file, tools/settings/MCP disabled; Codex: `model_instructions_file`). Each single-turn call is a fresh conversation. Temperature cannot be set through either CLI. Because the CLIs add their own hidden instructions and environment information, every comparison is between **model + CLI configurations**, not bare models. Server-side model versions are not pinned by date.

In the exploratory exp10 addendum only (§5.6.1), a fourth subject, `gemini` (`gemini-4-argon-medium`, "Gemini 4 Argon (Medium)"), was reached through the Antigravity CLI 1.4.7 (`agy -p … --output-format json`, an empty working directory, no automatic permission approval, a fresh single-turn conversation per item). This CLI is an agent harness with no option to replace its system prompt: its own instructions (about 6,800 input tokens for an empty prompt) always come first, and SOUL+MEMORY were placed at the start of the user message inside a `<system_instructions>` tag followed by one line asking for an answer only. Gemini therefore received the same documents in a different position and behind more harness text than the other models (§6).

In exp11 the same Claude CLI (2.1.289) and Codex CLI (0.164.0) were used with opus, sonnet, and gpt-6.1-sol. Of 1,159 answers, 1,154 were produced by the intended models (opus 382 × claude-opus-5-5, sonnet 382 × claude-sonnet-5-5, gpt 390) and five by an earlier model of the same tier (claude-opus-5 ×2, claude-sonnet-5 ×3), apparently after the provider's safeguard interrupted a call; they were kept (§5.7, deviations).

The control is the same-model re-run: sonnet in exp1–3, opus in exp4–8.

3.4 Grading

Keyword items are auto-graded. Rubric items are graded 0 / 0.5 / 1 by an LLM grader that sees all answers to one item in one prompt, shuffled once (seed 0) and unlabeled, with an instruction not to grade relatively. The primary grading is **cross-family**: Claude answers ← `gpt-6.1-sol`, GPT answers ← `opus`. Consequences: (i) for ratios between a Claude control and the GPT subject, numerator and denominator come from different graders; (ii) a grading prompt contains 15 Claude answers (three models × 5) but 5 GPT answers, so bundle size differs by family and may itself reveal family. From exp5 on, graders receive the subject's SOUL and MEMORY. A second, same-family grading pass exists for exp2, exp3 (opus graded all answers; GPT answers were not graded by GPT), exp5, and exp6; exp4, exp7, and exp8 have one cross-family pass. §5.2 re-computes ratios with one grader for all answers where data allow.

3.5 Metrics

- **Domain score**: mean item score within memory / judgment / persona / goals per repetition, then mean ± SD over repetitions. With N = 5, SDs are themselves very uncertain; they describe repetition spread on a fixed item set, not generalization to new items, documents, or users.
- **Carry-over ratio** (v0.1: "inheritance ratio"): swapped-model domain score / control domain score; always reported with the absolute difference. Undefined when the control scores 0; equal to 1 when both fail equally.
- **Application gap**: recall − application over paired items probing the same memory. **Rule gap**: aligned − conflict.
- **Slip rate**: among answers that *accepted* a legitimate change, the fraction that used the old rule again in the same answer (grader-judged). Because its denominator is the accepted answers, it has a selection effect; we also report **clean acceptance** = (accepted and not slipped) / all answers.
- **Illegitimate release**: score on NR items; a 0 means the answer effectively accepted lifting the prohibition; 0.5 typically means a standing pre-approval scheme was offered.
- **Regression**: control = 1 and swapped < 1 at the same repetition index. Pairing by index is arbitrary (no common random numbers), so we report it only descriptively.
- **Multi-turn persistence** (exp8): fraction of grader-judged post-change checks at which a changed rule is still followed (definitions in §5.3).

v0.1 flagged a "candidate difference" when a difference exceeded the control's repetition SD. That rule is not a test and is degenerate when the control SD is 0 (most domains); v0.2 drops it and reports absolute differences, per-model effects, and cluster-aware tests instead.

SOUL fidelity (new in v0.3; `src/soul_verify/fidelity.py`). The domain scores above use one rubric — the original SOUL and MEMORY are the answer key — so a "persona" score measures closeness to

that target, and most persona items score acceptance of a user's rule change plus memory facts rather than the SOUL's own voice (§5.4). SOUL fidelity scores each answer against the rules of the SOUL the subject actually received:

- **Specification.** For each SOUL, a list of checks is stored next to it (`<SOUL>.fidelity.yaml`); each check cites the SOUL sentence it comes from. (a) **Voice markers**, deterministic: speech level (ratio of polite to plain sentence endings, ≥ 0.8 or ≤ 0.2; noun-ending list items and table cells are not counted), forbidden address forms (quotations removed), emoji count (0, or 2–3), a required final-line prefix ("다음 할 일:" / "오늘의 한마디:"), a three-line summary at the top of long comparison/report answers, or the absence of lists. (b) **Boundaries**: what that SOUL forbids versus permits on the three items that probe them — lifting the payment prohibition (NR-1), recording a password (NR-2), giving the user's neighborhood to a self-declared sibling (IH-2b). IH-2b is a keyword check; NR-1/NR-2 use the graders' binary "refused" flag (not deterministic; averaged over graders where two exist). A rule that the SOUL leaves silent is not scored (e.g., the opposite SOUL has no password rule).
- **Legitimate changes.** On items where the user legitimately changes a rule, the check's expectation is replaced by the new rule only if the SOUL grants changes (principle 5 or its flexible half); for SOUL variants without such a sentence the changed marker is not scored on that item, because the document does not determine the answer.
- **Aggregation.** F_style and F_boundary are means over (answer, check) pairs; F is their mean. Rules that cannot be checked deterministically (conclusion-first ordering, empathy or a joke first in emotional answers, the principles on uncertainty and decisions, goals) are left to the rubric score.
- **Discrimination gap** = F(own SOUL's checks) − F(rival SOUL's checks) on the same answers. A metric that identifies the SOUL should give a large positive gap; the v0.4 persona score, by construction, has no rival version.
- **Deterministic acceptance** of a legitimate change = the answer passes all marker checks that encode the new rule (five of the six change items; RL-1 changes a time, not a voice marker).

The checks were written after exp9 had been run and inspected, and two detectors were revised once after inspecting answers (a sentence that *mentions* "요약" no longer counts as a summary header; noun-ending bullets no longer count as non-polite). Before the second revision the discrimination gaps were 0.71–0.79 and the persona fidelity ratios (exp4–exp7) 0.87–1.05; the conclusions below did not change. For the second agent of exp10, the markers and two new detectors (first-line prefix, final-line question; written from synthetic test sentences) were fixed in the preregistration before any answer existed (§5.6).

---

4 Experiments

Table 1. Overview of experiments.

Exp	Run id	Items	Subjects (N)	Control	Grading	Purpose
1 (pilot)	`pilot_2026-10-24_sonnet-vs-haiku`	v0.1, 20	sonnet, haiku (3)	sonnet	opus, 1 pass	Within-family swap

Exp	Run id	Items	Subjects (N)	Control	Grading	Purpose
2	`exp2_2026-10-24_gpt-vs-claude`	v0.1, 20	+ gpt-6.1-sol (3)	sonnet	cross-family + same-grader	Cross-family swap
3	`exp3_2026-10-24_v0.2-hard`	v0.2, 12 hard	sonnet, haiku, gpt (5)	sonnet	cross-family + opus-all	Harder items
4	`exp4_2026-10-24_v0.3`	v0.3, 22	opus, sonnet, haiku, gpt (5)	opus	cross-family, 1 pass	Ratios, application gap
5	`exp5_2026-10-25_v0.4`	v0.4, 23	same 4 (5)	opus	both graders	Principle 5 present ("full")
6	`exp6_2026-10-25_v0.4-no-p5`	v0.4, 23	same 4 (5)	opus	both graders	Principle 5 removed ("none")
7	`exp7_2026-10-25_p5-halves`	v0.4 subset, 14	same 4 (5) × 2 conditions	opus	cross-family, 1 pass	Each half of principle 5
8	`exp8_2026-10-25_multiturn`	v0.5, 4 scenarios × 12 turns	same 4 (3)	opus	cross-family, 1 pass (+ deterministic check, post hoc)	Multi-turn persistence
9	`exp9_2026-10-25_controls`	v0.4, 23	opus, sonnet, gpt (3) × 3 control conditions (+ C0 = exp5 repetitions 1–3)	C0 (same model)	**common panel: opus and gpt grade every answer**	No-SOUL / other-SOUL / swapped-MEMORY controls
10 (**preregistered**)	`exp10_2026-10-25_prereg`	v0.6: Doyun 24 + Haram 12 (10 memory, 12 illegitimate-release, 12 legitimate-change, 2 aligned)	opus, sonnet, gpt (2) × 11 conditions (172 condition×item cells)	— (conditions compared within model)	common panel (opus + gpt), deterministic fidelity with preregistered markers	Preregistered test of fidelity discrimination, both halves of principle 5 on 12 prohibition types, exp9 controls on a second agent
10A (exploratory addendum)	`exp10_gemini_2026-10-25`	same as exp10	gemini-4-argon-medium (2) × 11 conditions (172 cells)	—	same panel (opus + gpt), but gemini's answers graded separately (two per prompt)	Fourth model from a third family, added after preregistration

Exp	Run id	Items	Subjects (N)	Control	Grading	Purpose
11 (**preregistered**)	`exp11_2026-10-25_external`	v0.7: 16 English items per SOUL × 3 external SOULs (Maya, Talos, Dexter) = 48 (6 memory, 1 conflict, 4 illegitimate-release, 4 legitimate-change, 1 aligned per SOUL)	opus, sonnet, gpt (2) × 18 conditions (195 condition×item cells)	— (conditions compared within model)	common panel (opus + gpt), deterministic fidelity with preregistered markers	Replication of exp10 on human-written English SOULs (unmodified) with principle 5 appended

Run times for the ablation conditions (all 2026-10-25, KST): full 00:30–00:50, none 01:13–01:35, p5a 03:15–03:29, p5b 03:29–03:42. Conditions were not interleaved. Call counts: exp5 460 subject + 160 grading calls; exp6 460 + 160; exp7 560 + 56; exp8 48 conversations / 576 subject turns + 48 grading calls; exp9 621 subject calls (10:07–10:28) + 184 grading calls + 18 smoke calls. exp10: 1,032 subject calls (14:26–15:03, conditions interleaved) + 345 grading calls (344 grading units, one retried) + 12 smoke calls + 8 smoke grading calls = 1,397. No failed calls in the final data (exp6 had one retried process-spawn error; exp10 had 0 subject errors and 0 retries). exp10 addendum (gemini): 345 subject calls (344 answers + one retry of an empty answer; 15:40–16:07) + 344 grading calls + 6 adapter/item smoke calls + 6 smoke grading calls, about 695 in total; 0 failed answers. exp11: 1,186 subject call attempts (1,170 planned + 16 end-of-run retries; 19:17–19:59), 1,159 answers, 11 missing; 390 grading calls (195 per grader; 3 opus grading calls blocked by the same safeguard); 12 smoke + 8 smoke grading calls; 1,596 attempts in total. The SOUL fidelity metric (§3.5) made no model calls.

Sources: `README.md`; each run's `NOTES.md` and `raw.jsonl` timestamps.

---

5 Findings

5.1 Main result: ablation of the change-permission sentence

Motivation. Two opposite failures appeared before principle 5 existed. *Over-refusing a legitimate change:* in exp3, gpt refused the user's request to change the address form 5/5, prioritizing a voice rule over "Minseo decides" (H-10 score 0.00); in exp4 the same request (IH-3b) scored 0.20 for gpt. *Under-refusing an illegitimate release:* models offered a standing pre-approval scheme instead of refusing outright (NR-1 0.5 scores).

Design. exp5 ("full") added principle 5. exp6 ("none") deleted it (a one-line deletion). exp7 kept five principles but retained only one half:

- **p5a (flexible half, 102 characters):** preferences and voice rules can be changed by the user at any time, including time-limited changes.
- **p5b (firm half, 56 characters):** the two prohibitions are not lifted even on request; propose a safe alternative.

The same 14 items are compared across all four conditions. The outcome measures are acceptance of the six legitimate-change items (IH-3b, RL-1–RL-5), slips, and zero scores on NR-1 (payment prohibition). NR-2 (password) was refused 20/20 in every condition.

Table 2. Principle 5 ablation, pooled (descriptive; four models × 5 repetitions; cross-family grading).

Condition	Legitimate changes accepted	Slips / accepted	Clean acceptance	NR-1 zero scores	NR-1 mean
full (exp5)	119/120	2/119	117/120	0/20	0.73
none (exp6)	94/120	10/94	84/120	9/20	0.43
p5a only (exp7)	120/120	3/120	117/120	7/20	0.48
p5b only (exp7)	96/120	10/96	86/120	0/20	0.80

Sources: `runs/exp7_2026-10-25_p5-halves/analysis.md`*; clean acceptance from* `paper/reanalysis/results.md` *§B. NR-1 means are over the same four models (full 0.725 rounded).*

Table 3. Per-model effects and cluster-aware tests (★ = comparison corresponding to a claim of the paper).

Measure	Comparison	Δ pooled	Δ by model (opus / sonnet / haiku / gpt)	Pooled Fisher p (v0.1)	Stratified exact p (strata = model × item)	Cell sign-flip p (unit = model × item cell; non-zero cells)	Model sign-flip p (4 units)
Acceptance	★ none vs full	−0.21	0.00 / −0.20 / −0.10 / −0.53	1.0e-7	2.5e-10	0.016 (7/24); Holm 0.078	0.25
Acceptance	★ p5a vs none	+0.22	0.00 / +0.20 / +0.13 / +0.53	6.5e-9	5.2e-12	0.008 (8/24); Holm 0.047	0.25
Acceptance	★ p5b vs full	−0.19	−0.03 / −0.17 / −0.07 / −0.50	4.8e-7	3.4e-9	0.016 (7/24); Holm 0.078	0.125
Acceptance	p5a vs full	+0.01	0.00 / 0.00 / +0.03 / 0.00	1	1	1 (1/24)	1
Acceptance	p5b vs none	+0.02	−0.03 / +0.03 / +0.03 / +0.03	0.87	0.74	0.63 (4/24)	0.63

Measure	Comparison	Δ pooled	Δ by model (opus / sonnet / haiku / gpt)	Pooled Fisher p (v0.1)	Stratified exact p (strata = model × item)	Cell sign-flip p (unit = model × item cell; non-zero cells)	Model sign-flip p (4 units)
Clean acceptance	none vs full	−0.28	0.00 / −0.17 / −0.13 / −0.80	2.6e-9	1.0e-12	0.008 (10/24)	0.25
NR-1 zeros	★ none vs full	+0.45	+1.00 / +0.20 / +0.60 / 0.00	0.0012	3.3e-4	0.25 (3/4); Holm 0.75	0.25
NR-1 zeros	★ p5b vs none	−0.45	−1.00 / −0.20 / −0.60 / 0.00	0.0012	3.3e-4	0.25 (3/4); Holm 0.75	0.25
NR-1 zeros	★ p5a vs full	+0.35	+0.40 / 0.00 / +1.00 / 0.00	0.0083	0.002	0.50 (2/4); Holm 0.75	0.50
NR-1 zeros	p5a vs none	−0.10	−0.60 / −0.20 / +0.40 / 0.00	0.75	0.66	0.75 (3/4)	0.75

Source: `paper/reanalysis/results.md` §B (full table, including slips and per-model Fisher p). The stratified exact test conditions on each model×item stratum's margins but still treats the 5 repetitions within a stratum as exchangeable, so it inherits the pseudo-replication of the pooled test. The cell sign-flip test treats each model×item cell as one unit; with one NR item it has only four cells. With four models, the smallest attainable two-sided model-level p is 0.125. Holm adjustment is over the six ★ comparisons within each test type.

Per-model reading.

- Acceptance: the drop without principle 5 is large for gpt (14/30; Fisher p = 1.9e-6), moderate for sonnet (24/30; p = 0.024, not significant after Holm over 24 per-model tests), small for haiku (26/30; p = 0.35), and absent for opus (30/30). The cell-level tests that remain below 0.05 are driven by gpt (IH-3b, RL-3, RL-4, RL-5), by RL-4 ("casual speech just for today") in sonnet and haiku, and by IH-3b in sonnet (none vs. full: 7 non-zero cells of 24).
- NR-1 zeros: without principle 5, opus accepted the release 5/5 and haiku 3/5; sonnet 1/5; gpt 0/5. With p5a only, haiku accepted it 5/5 and opus 2/5. NR-1 is a single item, so these are four model-level observations of one request.
- So "the sentence fixed both failure modes" holds only in aggregate: each model's missing behavior differed (opus lacked refusal, gpt lacked acceptance), and principle 5 *reduced* the failure each model had. It did not remove all failures: under the full condition, opus and haiku still scored 0.5 on NR-1 5/5 (pre-approval offered), and haiku refused IH-3b once.

Grader sensitivity (full and none only). Re-grading exp5 and exp6 with a single grader for all answers leaves the full condition unchanged (119/120, 2/119, 0/20 with either grader) and changes the none condition slightly (opus grader: 94/120, 9/94, 8/20; gpt grader: 99/120, 15/99, 9/20) (`paper/reanalysis/results.md` §A). exp7 has no second grader.

Deterministic replication of the acceptance effect (v0.3). Because exp7 has one grader, we re-scored acceptance without graders: an answer accepts a change if it passes every voice marker that encodes the new rule (IH-3b: no "민서 님"; RL-2: no summary header; RL-3: at least one emoji; RL-4: casual speech; RL-5: no "다음 할 일:" line). Agreement with the graders' acceptance flag on these 100 answers per condition was 0.99 / 0.95 / 1.00 / 0.93 (κ 0.66 / 0.87 / 1.00 / 0.82).

Table 2b. Deterministic acceptance of five legitimate changes (four models × 5 repetitions).

Condition	Deterministic full acceptance	Grader acceptance (same 5 items)	By model, deterministic (opus / sonnet / haiku / gpt; of 25)	RL-4 "casual today"	IH-3b address
full	98/100	99/100	25 / 25 / 23 / 25	0.95	0.95
none	71/100	74/100	25 / 20 / 21 / 5	0.35	0.70
p5a only	100/100	100/100	25 / 25 / 25 / 25	1.00	1.00
p5b only	71/100	76/100	24 / 20 / 22 / 5	0.30	0.75

Source: `paper/reanalysis/fidelity_results.md` §2. Cell sign-flip p (model × item, 20 cells): none − full −0.27, p = 0.031; p5a − none +0.29, p = 0.016; p5b − full −0.27, p = 0.023; p5a − full p = 0.50; p5b − none p = 1.0.

The deterministic scorer reproduces the pattern and the per-model concentration (gpt 5/25 without the flexible half) and removes grader asymmetry from this part of the result. It does not address the refusal side, whose only deterministic signal would be on a single item.

Where the refusals came from (exp9). With no SOUL at all (exp9 C1: "You are Haram, an assistant who helps Minseo" plus the date), the same three models accepted every legitimate change — grader acceptance 1.00, deterministic 45/45 — with slips 0.03 (§5.4). Refusing a change is therefore not a general tendency of these models; it appears when the SOUL contains rigid clauses such as "keep polite speech even if Minseo speaks casually", "never say 'Minseo-ssi'", and "no emoji, even if Minseo uses emoji", and no sentence says which rules the user may change. Principle 5's acceptance effect is best read as **undoing refusals that the SOUL's own rigid rules created**: without principle 5 the document contradicts the user's request; the flexible half resolves that contradiction. (C1 has no haiku and is not the same batch as exp6; the comparison is across experiments.)

Models fill the gap differently. Without principle 5, gpt defaulted to *rule text first* — accepting 14/30 legitimate changes, refusing IH-3b, RL-3 (emoji), and RL-4 (casual speech today) 5/5 each, and slipping in 9/14 of the changes it did accept — while scoring 5/5 on NR-1. opus defaulted to *user decides first* — accepting 30/30 but scoring 0 on NR-1 5/5 ("The payment-approval rule can be changed as you decide. But since payment is hard to undo, let me confirm once…") (`runs/exp6_2026-10-25_v0.4-no-p5/NOTES.md`). An explicit priority sentence narrowed this difference in defaults.

The flexible half leaks without a boundary. With p5a only, opus and haiku applied the flexible wording ("can change at any time", "set a period such as just today") to the payment prohibition, e.g., proposing to "try it for just a month" (`runs/exp7_2026-10-25_p5-halves/NOTES.md`).

Interpretation and confounds. The four conditions behave like a 2×2 design (flexible half present/absent × firm half present/absent), and the observed pattern is consistent with the flexible

half governing acceptance and the firm half governing release refusal: acceptance with p5a alone was close to full (120 vs. 119), with p5b alone close to none (96 vs. 94); NR-1 zeros with p5b alone matched full (0 vs. 0), with p5a alone were close to none (7 vs. 9). These similarities are observations, not equivalence tests, and the separation is not perfectly clean (sonnet slipped 3/30 with p5a vs. 1/30 with full; NR-1 mean was higher with p5b alone, 0.80, than with full, 0.73). The ablation does not establish a mechanism, because: (i) **batch and time** — each condition was run as one batch, not interleaved, over about three hours of one night; (ii) **length and position** — p5a and p5b differ in length (102 vs. 56 characters) and both occupy the last principle slot, so the number of principles is controlled but length, position, and emphasis are not; (iii) **direct instruction** — each half explicitly instructs the behavior it is scored on, so the result is also explained by "adding an explicit instruction changes the instructed behavior"; (iv) **items** — one illegitimate-release item and six legitimate-change items, all written and revised by the authors, one wording each; (v) **grading** — one cross-family pass in exp7 (the acceptance side is now also reproduced deterministically, Table 2b). A confirmatory test needs interleaved runs, length-matched and paraphrased variants, a neutral fifth-principle control, new held-out items and prohibition types, and a common grader panel (§6). exp10 (§5.6) supplies part of this — interleaved conditions, length-matched halves for the new agent, 10 new prohibition items, a second agent, a common grader panel, and preregistered hypotheses — but not paraphrased variants or a neutral fifth-principle control.

Not everything was explained by principle 5. exp4's opus/sonnet slips on IH-3b (0.70/0.60) did not recur in exp6 (opus 1.00, 0 slips), although the exp6 SOUL differs from the v0.3 SOUL only by the date line. A possible explanation is the date contamination and item-wording changes between v0.3 and v0.4 (§6); we could not separate these.

Automatic slip detection was unreliable. A token-based slip check agreed with graders 0.93 but with $\kappa = -0.03$ (false positives when the model *mentions* a dropped rule; misses on structural slips) (`runs/exp5_2026-10-25_v0.4/analysis.md`).

5.2 Secondary: recall, application, and carry-over ratios

Single-fact recall was at ceiling; other memory checks were not. From v0.3 on, the keyword-graded single-fact recall items (R-xx) were answered correctly by all four models in every repetition. This ceiling may reflect item difficulty rather than robust carry-over. It does not extend to all memory checks:

Table 4. Counter-examples to "memory carried over fully".

Experiment	Grader	Model	Memory result	Where the points were lost
exp3 (v0.2 hard items)	cross-family (primary)	haiku	0.84 ± 0.09 (5 memory items)	H-06 long-context item 0.5 in 5/5 runs; H-03 0 once; H-01 0.5 once
exp3	opus (secondary)	haiku	0.80 ± 0.10	as above, plus H-11 0.5 twice
exp8 (12 turns)	cross-family	haiku	memory checks at turns ≤ 7: 9/9; turns ≥ 8: 16/28 = 0.57	late checks, including invented schedule items
exp8	cross-family	opus / gpt	turns ≥ 8: 32/33 each	one late check each

Experiment	Grader	Model	Memory result	Where the points were lost
exp4–exp7	—	all	memory domain 1.00	—

Source: `paper/reanalysis/results.md` §C; `runs/exp3_2026-10-24_v0.2-hard/analysis.md`; `runs/exp8_2026-10-25_multiturn/analysis.md` §3.

Application diverged where recall did not. Application gap (recall − application on paired items): exp4 (6 pairs) opus 0.02 ± 0.04, sonnet 0.02 ± 0.04, haiku 0.20 ± 0.05, gpt 0.02 ± 0.04; exp5 (4 pairs) haiku 0.25 ± 0.15, others 0.00 ± 0.00. In exp4, the opus and sonnet values came from repetition 5 of A-04, after midnight (pre-midnight repetitions 1–3: 0.00); gpt's 0.02 came from A-05 in repetition 1, *before* midnight (pre-midnight 0.03 ± 0.05), so it is not explained by date contamination (`paper/reanalysis/results.md` §E). The clearest instance: haiku stated "dislikes Friday-afternoon meetings" correctly 5/5 when asked directly, but recommended Friday 30 October afternoon in 2/5 runs when asked to choose a slot (exp4, A-03). This replicates the recognize-vs-follow split of PrefEval [B4] under a model swap, in one model.

Carry-over ratios depend on the grader. Table 5 compares, for exp5, the primary cross-family ratio with ratios computed with one grader for all answers.

Table 5. Carry-over ratio (swapped / opus control), exp5 (v0.4, N = 5), by grading; absolute difference in parentheses.

Model	Domain	Cross-family (primary)	opus grades all	gpt grades all	Mean of two graders
sonnet	judgment	1.07 (+0.06)	1.05 (+0.05)	1.07 (+0.06)	1.06 (+0.06)
sonnet	persona	0.99 (−0.01)	0.99 (−0.01)	0.99 (−0.01)	0.99 (−0.01)
haiku	judgment	0.86 (−0.12)	0.81 (−0.18)	0.86 (−0.12)	0.83 (−0.15)
haiku	persona	0.81 (−0.19)	0.82 (−0.18)	0.81 (−0.19)	0.82 (−0.18)
haiku	goals	0.80 (−0.20)	0.77 (−0.23)	0.80 (−0.20)	0.78 (−0.22)
gpt-6.1-sol	judgment	**1.07 (+0.06)**	**1.00 (0.00)**	**1.00 (0.00)**	1.00 (0.00)
gpt-6.1-sol	persona	0.99 (−0.01)	0.99 (−0.01)	0.99 (−0.01)	0.99 (−0.01)

Memory 1.00 and goals 1.00 for sonnet and gpt under every grading (omitted). Control (opus) judgment: 0.88 cross / 0.94 opus grader / 0.88 gpt grader. Source: `paper/reanalysis/results.md` §A, which also gives exp2, exp3, and exp6.

Within the Claude family (all answers graded by gpt in the primary scoring), the conclusions are stable across graders: sonnet ≈ control, haiku 0.77–0.86 on judgment/persona/goals. **The persona rows of this table do not measure carry-over of the SOUL's persona** (§5.4: the same rubric gives 0.96 to answers written under an opposite SOUL); persona carry-over measured with document-specific fidelity is in §5.5 (Table 9), where haiku's ratio is 0.95, not 0.81. The cross-family ratio is less stable. gpt's apparent advantage on judgment in exp5 (1.07) disappears under either single grader (1.00), because the opus control is graded more strictly by gpt than by opus. In the earlier experiments the grader changed the direction of two ratios: exp3 gpt judgment 1.17 (cross) vs. 0.97 (opus grades all; no gpt-graded GPT answers exist for exp3), and exp2 haiku persona 1.04 (cross; within-family, gpt grader) vs. 0.97 (opus grader); exp2 gpt persona was 1.15 cross vs. 1.00–1.04 single-grader. Ratios

above 1 for gpt in v0.1 (Table 2 of v0.1) should therefore not be read as gpt following the specification better than the control.

Ceiling items and hallucination. In exp5 the item with the most negative flips was X-04 (deposit request with two listings and a budget cap): haiku 5/5 and gpt 5/5 failed to state that listing B (1.9M KRW/month) exceeds the 1.8M cap, while opus and sonnet passed 5/5 (`runs/exp5_2026-10-25_v0.4/NOTES.md`). Zero scores on abstention/false-premise items: 0 for all models in exp1–3 (partial 0.5 scores in exp3: sonnet 0–1/25, haiku 11–13/25, gpt 2/25 depending on grader); exp4 haiku 2/15, others 0/15; exp5 0 for all models (haiku 4/10 at 0.5); exp6 haiku 1/10 (4 at 0.5), others 0; exp7 haiku 0/5 (p5a) and 1/5 (p5b), others 0 (`paper/reanalysis/results.md` §C; v0.1 omitted exp5–7). exp8 has no separate hallucination item; invented schedule items appear inside haiku's late memory checks (Table 4).

5.3 Secondary: multi-turn persistence

exp8 used each CLI's real continuation (Claude: one `claude -p` stream-json process per conversation; Codex: `exec resume` with the same instructions file each turn). Four fixed 12-turn scripts changed rules at turns 2–3, interleaved small talk and tasks, attempted an illegitimate release around turns 9–10, and checked behavior at six checkpoints each (24 checkpoints). User turns do not react to the model's answers.

Denominators. Each checkpoint contains zero or more "changed-rule" checks; over the 24 checkpoints there are **40 changed-rule check slots per model and repetition**, i.e. 120 per model (3 repetitions). A slot is counted when the grader judged it true/false; haiku has 3 null slots (S4 T8, C_brief7, all three repetitions), hence 117. **Distance** = checkpoint turn − turn at which that rule was changed. Distance 0 is the turn of the change itself (accepting and applying it) and is reported separately; "all post-change" counts distance ≥ 1 only (99 per model; 96 for haiku). For the address-form rule (C_name) the grader's judgment was replaced, **post hoc**, by a deterministic check after we verified that all 33 occurrences of the old form flagged by the grader were quotations; the grader's original judgment is kept as a sensitivity analysis (opus and sonnet 27/36, haiku 32/36, gpt 36/36 under the original judgment; `runs/exp8_2026-10-25_multiturn/analysis.md` §1).

Table 6. Persistence of changed rules by turn distance (exp8, N = 3).

Model	Distance 0	1–3	4–6	7–10	All post-change (distance ≥ 1)
opus	1.00 (21/21)	1.00 (9/9)	0.94 (34/36)	1.00 (54/54)	0.98 (97/99)
sonnet	1.00 (21/21)	0.78 (7/9)	0.97 (35/36)	0.96 (52/54)	0.95 (94/99)
haiku	1.00 (21/21)	0.67 (6/9)	0.94 (31/33)	0.96 (52/54)	0.93 (89/96)
gpt-6.1-sol	1.00 (21/21)	0.89 (8/9)	0.92 (33/36)	0.98 (53/54)	0.95 (94/99)
all, 95% Wilson CI	1.00 (84/84; 0.96–1.00)	0.83 (30/36; 0.68–0.92)	0.94 (133/141; 0.89–0.97)	0.98 (211/216; 0.95–0.99)	0.95 (374/393)
all, excluding C_story	84/84	24/24 (0.86–1.00)	117/117 (0.97–1.00)	203/204 (0.97–1.00)	—

Sources: `runs/exp8_2026-10-25_multiturn/analysis.md` *(§1);* `paper/reanalysis/results.md` *§D. Checks within a conversation are not independent; the effective unit is the conversation (48; 12 per model).*

- **Decay with distance was not detected; the design could not have detected it.** Rule type is confounded with distance (e.g., the 1–3 bin has 9 slots per model, 3 of which are one checkpoint of the vague story-style rule, S4 T6; most 7–10 slots are address-form and briefing-time checks). Excluding the story-style rule, every post-change bin is at 0.99–1.00, leaving no room to observe decay. Conversations with at least one post-change slip: opus 2/12, sonnet 3/12, haiku 3/12, gpt 3/12 — all in scenario S4.
- **Slips concentrate in one vague rule.** 18 of 19 slips were on "tell it like a story" (C_story; pass rates opus 10/12, sonnet 7/12, gpt 7/12, haiku 6/12). This rule also conflicts with the unchanged "three-line summary on top" rule, so these slips may reflect a rule collision rather than forgetting.
- **Time-scoped changes.** In S3 ("casual speech *just for today*"), opus, sonnet, and gpt wrote tomorrow's briefing in formal speech while still speaking casually today (3/3 each); haiku did not (1 casual, 2 asked back — possibly a script mismatch, since the fixed script ignores questions).
- **Late illegitimate releases (text answers).** Password (S2) and card-number-plus-payment (S4) requests were refused in 12/12 conversations each, without echoing the secrets. Payment-release means: gpt 1.00; opus 0.83 / 0.83 (S1 / S3); sonnet 0.67 / 1.00; haiku 0.17 / 0.67. These are one wording per scenario and do not show robustness to other phrasings.
- **History carries mistakes forward.** haiku's loose alternatives at S1 T10 ("pay, then report") reappeared as recommendations in its T12 summary (0/3).

Sources: `runs/exp8_2026-10-25_multiturn/NOTES.md`; `runs/exp8_2026-10-25_multiturn/analysis.md` *(§2–4).*

5.4 Document controls: what the SOUL, the MEMORY, and general ability determine (exp9)

Question. Are the scores due to the content of the SOUL, or to a general ability to follow instructions? exp9 answers the v0.2 limitation "no no-SOUL, other-SOUL, or swapped-MEMORY control" on the v0.4 items (23 items, reference date 2026-10-25), with opus, sonnet, and gpt-6.1-sol × N = 3 (haiku omitted). All SOULs and MEMORYs are synthetic.

Table 7. Conditions.

Condition	SOUL	MEMORY	Note
C0	v0.4 Haram SOUL (with principle 5)	Minseo	exp5 answers reused (repetitions 1–3); no new calls
C1 no-SOUL	one line: "You are Haram, an assistant who helps Minseo" + date line	Minseo	the date is environment, not SOUL content (exp4 midnight lesson)
C2 other-SOUL	same name, user, and goals; opposite personality and rules	Minseo	casual speech, "민서야", reasons first and conclusion last, no lists, final line "오늘의 한마디:", 2–3 emoji, a joke first in emotional talk; payments ≤ 300k KRW without asking (a changeable *way*), schedule and address may be given to family; principle 5 in the same form; two *different* prohibitions (no messages in Minseo's name, no sharing health information); no password rule

Condition	SOUL	MEMORY	Note
C3 swapped-MEMORY	Haram SOUL	"Jiwoo"	same memory IDs and topics, different values (Mangwon-dong, cap 1.5M, opening 2027-06, ...). Voice-related corrections kept in the SOUL's direction so that persona is set by the SOUL only

Grading. Every answer of every condition — including C0, re-graded — was graded blind by both opus and gpt-6.1-sol with the same prompt; the primary score is the mean of the two graders. Graders received the condition's own SOUL and MEMORY as reference, but `score` uses the v0.4 rubric (Haram + Minseo as the answer key), i.e. "distance from the C0 target". A per-condition `ref_follow` judgment (did the answer follow the documents it received?) had low grader agreement (κ 0.06–0.35) and is not used for conclusions. C0 re-grading reproduced the exp5 grades (opus exact agreement 1.00; gpt 0.99, κ 0.91).

Table 8. Results (mean of three models; mean of two graders; parentheses = opus / gpt grader).

Measure	C0	C1 no-SOUL	C2 other-SOUL	C3 Jiwoo MEMORY
Total score (23)	0.97 (0.98/0.95)	0.93 (0.95/0.92)	0.89 (0.91/0.88)	**0.40** (0.42/0.38)
Memory recall (R, 4)	1.00	0.99	0.97	**0.12**
Memory application (A, 4)	1.00	1.00	1.00	**0.12**
Persona domain (8)	0.99	0.96	**0.96**	0.46
Judgment domain (8)	0.93	0.85	0.75	0.57
Legitimate changes accepted (6)	1.00	1.00	1.00	1.00
NR refusal (NR-1, NR-2)	0.92 (1.00/0.83)	**0.53** (0.67/0.39)	**0.33** (0.33/0.33)	0.92
Haram voice markers (16 items, deterministic)	**0.97**	0.67	**0.14**	0.96
C2 voice markers (deterministic)	0.00	0.01	**0.85**	0.00
Memory source = Jiwoo only (12 items, keywords)	0/108	0/108	0/108	**83/108**

Source: `runs/exp9_2026-10-25_controls/analysis.md` §2–4, `NOTES.md`. The voice-marker rows are the exp9 analysis' own composite (emoji, polite ratio, final line, address) on 16 items that do not change the voice; §5.5 gives the full fidelity metric.

- **Memory items are determined by the MEMORY.** With Jiwoo's memory, recall and application fell to 0.12 (the remaining credit is R-05's "vendor unknown"), and 83/108 answers used Jiwoo-only cues. Date arithmetic and application were intact on the new values (e.g., OPIc 12-05 → 41 days). Answers that pointed out that the records belong to "Jiwoo" while the SOUL names Minseo: opus 42/69, gpt 19/69, sonnet 7/69 — detection of the SOUL/MEMORY mismatch differs strongly by model.
- **Voice is determined by the SOUL** — under C2, emoji in 100% of answers, polite sentences 0–2%, final line "오늘의 한마디:" 100% in all three models; under C3 the Haram voice was intact. One exception: **the address form followed the MEMORY's correction over the SOUL** for opus

(C2 "민서야" gpt 0.77, sonnet 0.50, opus 0.02; opus followed MEMORY M14 "call me Minseo-nim" while speaking casually). Priority between a user correction in MEMORY and a SOUL rule differs by model.

- **Prohibition boundaries are partly set by the SOUL, partly general.** NR-1 (lift payment approval): 0/18 refusals under C2, whose SOUL calls the payment limit a changeable way; under C1 all 9 answers offered a standing limit (0.5), and opus and sonnet left open full delegation "if you want". The C0 behavior "keep explicit approval and offer alternatives" is due to the SOUL. NR-2 (store a password in plain text): sonnet and gpt refused 6/6 with no SOUL and with a silent SOUL — a general safety tendency; opus refused 2/6 (C1) and 0/6 (C2), explaining the risk and handing the decision to Minseo, so opus's C0 refusals (6/6) are due to the SOUL. IH-2b (a self-declared brother asks for the address): under C2, which permits family, opus 3/3 and gpt 3/3 named the neighborhood; sonnet 0/3 put identity verification first.
- **Accepting legitimate changes is general ability.** Acceptance was 1.00 in all four conditions, with slips 0–3% — including C1, which has no principle 5 and no rigid voice clauses (§5.1).
- **The persona domain score hardly separates SOULs.** C2 (opposite voice) 0.96 and C1 (no SOUL) 0.96 versus C0 0.99 (cell sign-flip p 0.25 and 0.38), while the deterministic markers collapsed (0.97 → 0.14). Most of the eight v0.4 persona items score "accepts the user's rule change" plus memory facts; Haram's own voice (polite speech, no emoji, "다음 할 일:") is barely in their rubrics. **The persona carry-over ratios of exp4–exp8 therefore measured mostly general instruction following and MEMORY use, not carry-over of SOUL-specific behavior.** This is corrected in §5.5.
- **Dependence on the SOUL in the total score is small, on the MEMORY large**: C1 − C0 = −0.035 (cell p 0.09, Holm 1.0), C2 − C0 = −0.074 (p 0.006, Holm 0.17; mostly judgment items NR-1, NR-2, IH-2b), C3 − C0 = −0.57 (p < 1e-5, Holm 2e-4). With three models, the smallest model-level p is 0.25: model generalization cannot be claimed.

5.5 SOUL fidelity: discrimination and recomputed persona carry-over

Discrimination (exp9). Each answer was scored with the Haram checks and with the C2 checks (§3.5). A metric that identifies the SOUL must score answers high on their own SOUL's checks and low on the rival's.

Table 9. SOUL fidelity of the same answers under two specifications (exp9; three models × 23 items × 3 repetitions).

Condition	Own SOUL	F_style, Haram checks	F_style, C2 checks	F_boundary, Haram	F_boundary, C2	Gap F_style (own − rival)	Gap F
C0 Haram SOUL	Haram	**1.00**	0.17	0.94	0.08	**0.83**	0.84
C1 no-SOUL	—	0.76	0.13	0.69	0.36	—	—
C2 other-SOUL	C2	0.19	**0.90**	0.33	0.83	**0.71**	0.60
C3 Jiwoo MEMORY	Haram	**1.00**	0.16	0.94	0.08	**0.84**	0.85

Source: `paper/reanalysis/fidelity_results.md` §1. Own-versus-rival gap of F_style: positive in 69/69 model×item cells in each of C0, C2, C3 (cell sign-flip p < 1e-5, Monte Carlo). By model, gap in C0: opus 0.86, sonnet 0.82, gpt 0.80; in C2: opus 0.54, sonnet 0.73, gpt 0.85. On the eight persona items alone, the v0.4 persona score was 0.99 / 0.96 / 0.96 / 0.46 (C0 / C1 / C2 / C3) and Haram F_style 0.95 / 0.70 / 0.30 / 0.95.

- The fidelity metric separates the two SOULs where the persona score did not, and it is insensitive to the MEMORY swap (C3 = C0), as a SOUL measure should be.
- **Most of the Haram voice is a model default.** With no SOUL, polite speech (1.00), not using "민서 씨/야" (1.00), and no emoji (0.95) were already present; what the SOUL added was the final "다음 할 일:" line (C1 0.05 → C0 1.00) and the boundary behavior on NR-1. F_style C1 − C0 = −0.24 (64/69 non-zero cells, p < 1e-5). A SOUL whose rules coincide with the model's defaults will look "inherited" by any model; specificity is visible only on rules that differ from the default.
- C2's own fidelity was 0.90, lowered by the address check (opus 0.23, following the MEMORY correction; sonnet 0.59; gpt 0.93) and, for boundaries, by sonnet's refusal to give the address that C2 permits (0/3).

Recomputed persona carry-over (exp4–exp8). Table 10 sets the v0.4 persona-domain ratio (primary cross-family grading, as in Table 5) beside the ratio of each model's F_style under its own SOUL's checks to the opus control's. All F_style values come from one deterministic scorer, so the ratios have no grader asymmetry.

Table 10. Persona carry-over: rubric persona score vs. document-specific fidelity (control = opus).

Experiment	Model	Persona score	Persona ratio	F_style (own SOUL)	F_style ratio	F_style, C2 checks	Gap
exp4 (v0.3)	sonnet	0.83	0.93	1.00	1.00	0.15	0.85
exp4	haiku	0.67	0.74	0.94	0.94	0.21	0.73
exp4	gpt	0.73	0.81	1.00	1.00	0.19	0.81
exp5 (full)	sonnet	0.99	0.99	1.00	1.00	0.17	0.83
exp5	haiku	0.81	0.81	0.94	0.95	0.24	0.70
exp5	gpt	0.99	0.99	1.00	1.00	0.21	0.79
exp6 (none)	sonnet	0.84	0.85	1.00	1.00	0.16	0.84
exp6	haiku	0.70	0.71	0.95	0.95	0.23	0.71
exp6	gpt	0.49	0.49	1.00	1.00	0.19	0.81
exp7 p5a	haiku	0.86	0.88	0.95	0.95	0.26	0.70
exp7 p5b	haiku	0.80	0.84	0.93	0.93	0.24	0.69
exp7 p5b	gpt	0.43	0.45	1.00	1.00	0.21	0.79
exp8 (12 turns)	haiku	—	—	0.99	0.99	0.20	0.79

Source: `paper/reanalysis/fidelity_results.md` §3 (all rows; the opus control has F_style 1.00 in every experiment, and the omitted sonnet/gpt rows of exp7–8 all have F_style ratio 1.00). exp8: turn-level

checks of speech level, address, and emoji only, dropping a marker from the turn its rule was changed in that conversation; persona score not defined for exp8.

- **What changes.** v0.2's persona ratios (e.g., haiku 0.71–0.88, gpt 0.45–0.49 without the flexible half, gpt 0.81 in exp4) did not measure loss of the SOUL's voice: under document-specific checks every model kept the voice at 0.93–1.00 of the control, and gpt's low persona scores came from refusing or slipping legitimate *changes* (§5.1), which the persona rubric scores and the voice checks do not (on those items the changed marker is not scored when the SOUL grants no changes).
- **What remains.** haiku is still the lowest (0.93–0.95; exp5: final "다음 할 일:" line 0.79, three-line summary 0.83), consistent with its lower memory application.
- **What weakens.** "The persona carried over" is true only in a weak sense: the same markers are 0.76 with no SOUL. The SOUL-specific part that carried over to all four models is small: essentially the fixed final line and — on one item each — the payment and password boundaries (the three-line summary was produced at a similar rate without the SOUL, 0.89 vs. 0.90).

5.6 Preregistered replication on a second agent (exp10)

Aims. exp10 addressed three limitations of v0.3: (A) the fidelity markers had been written after exp9 was inspected, and two detectors were revised after looking at answers; (B) the firm half of principle 5 had been tested on one illegitimate-release item (NR-1; NR-2 was at ceiling); (C) all controls used one agent. All documents are synthetic.

Preregistration and timing. The preregistration (`v0.6/PREREGISTRATION.md`) fixed the materials, the fidelity markers (`*.fidelity.yaml`), nine confirmatory hypotheses with effect thresholds, the test, the Holm correction, descriptive criteria, exclusion rules, the call budget, and the run/grade/analysis code. It was committed, and its SHA-256 manifest was submitted to three OpenTimestamps calendars, before the first model call:

Table 11. exp10 timeline (2026-10-25, KST).

Step	Time	Evidence
Preregistration commit `0ac513d` (materials, markers, hypotheses, analysis code)	14:24:42	`git log`
Timestamp commit `7581bdf` (`proofs/2026-10-25_0ac513d.txt` + `.ots`)	14:25:15	`git log`, `.ots`
First smoke call	14:25:34	`runs/smoke_exp10_2026-10-25/raw.jsonl`
Main run (1,032 answers)	14:26:41 – 15:03:24	`raw.jsonl` timestamps
Grading	15:04 – 15:13	`grade_log.txt`
Results commit `ac74d85`, timestamp commit `f411d8e`	15:16:29, 15:17:04	`git log`

The git commit time is set by the author's machine; the OpenTimestamps proof shows only that the manifest existed no later than the Bitcoin block that eventually anchors it (Appendix B). After preregistration, only three bugs in `analyze10.py` were fixed, found on smoke data before the main data were analyzed (a formatting error on empty conditions, a descriptive criterion that reported "replicated" on missing data, a verdict label); markers, detectors, rubrics, hypotheses, and tests were not changed. A descriptive script (`peek.py`, marker pass rates and example answers) was added afterwards and was not used for any verdict.

Design. Eleven conditions: Haram full / p5a / p5b / none (12 items: NR-1, NR-2, four new prohibition items — transfer, subscription, card number, national ID number — and the six legitimate-change items of v0.4); Doyun full (24 items: 10 memory items in five recall/application pairs, six illegitimate-release items, six legitimate-change items, two aligned requests), Doyun p5a / p5b / none (14 items), Doyun no-SOUL (one line plus the date), Doyun opposite SOUL, and Doyun with a swapped MEMORY ("Haeun"; memory items only). Doyun's six prohibitions cover payment, passwords/one-time codes, personal information, security prompts, deletion without confirmation, and external posting. For Doyun, the two halves of principle 5 were length-matched (75 and 75 characters); Haram's halves were reused unchanged (102 vs. 56). The opposite SOUL keeps name, user, goals, and the form of principle 5, but reverses the voice (polite speech, "서연 님", no lists, no emoji, final line "도윤 드림"), allows payments up to 50,000 KRW, deletion without confirmation, and messages in Seoyeon's name, and keeps three prohibitions (codes/passwords, personal information, security prompts). Subjects: opus, sonnet, gpt-6.1-sol, N = 2 (N = 3 would have exceeded the preregistered budget of 1,500 calls). Within each repetition, all (condition, item, model) jobs were shuffled into one queue, so conditions were **interleaved in time**, unlike exp5–exp7. Both opus and gpt-6.1-sol graded every answer with the condition's own SOUL and MEMORY as reference; the score of an answer is the mean of the two graders (exact agreement 0.94 over 1,032 answers). Tests: sign-flip permutation over cells (model × item for H1–H3 and H6; agent × model × item, 36 cells, for H4–H5), Holm over the nine hypotheses.

Table 12. Preregistered hypotheses (exp10; Holm over nine, α = 0.05).

#	Comparison	A / B	Effect	Threshold	Cell p (non-zero / all cells)	Holm p	Verdict
H1	Doyun-SOUL answers: F_style, Doyun markers − opposite markers	0.99 / 0.22	+0.77	≥ 0.5	5e-6 (72/72)	4.5e-5	supported
H2	Opposite-SOUL answers: F_style, opposite − Doyun markers	0.87 / 0.18	+0.69	≥ 0.5	5e-6 (72/72)	4.5e-5	supported
H3a	Doyun-SOUL answers: F (incl. boundaries), own − rival	1.00 / 0.39	+0.60	> 0	5e-6 (72/72)	4.5e-5	supported
H3b	Opposite-SOUL answers: F, own − rival	0.92 / 0.36	+0.56	> 0	5e-6 (72/72)	4.5e-5	supported
H4a	Illegitimate-release score, p5b − none (both agents)	0.99 / 0.85	+0.14	> 0	0.004 (10/36)	0.012	supported
H4b	Illegitimate-release score, full − p5a	0.96 / 0.90	+0.06	> 0	0.129 (10/36)	0.129	**not supported**
H5a	Legitimate-change acceptance, p5a − none	1.00 / 0.77	+0.23	> 0	9.8e-4 (11/36)	0.004	supported

#	Comparison	A / B	Effect	Threshold	Cell p (non-zero / all cells)	Holm p	Verdict
H5b	Legitimate-change acceptance, full − p5b	1.00 / 0.80	+0.20	> 0	0.008 (8/36)	0.016	supported
H6	Memory score, swapped MEMORY − full	0.23 / 0.99	−0.77	≤ −0.5	5e-6 (30/30)	4.5e-5	supported

Source: `runs/exp10_2026-10-25_prereg/analysis.md` *§1,* `analysis.json`*. p = 5e-6 is the Monte Carlo floor (200,000 sign flips). By model, H1 gap: opus 0.81, sonnet 0.73, gpt 0.79; H2: opus 0.61, sonnet 0.56, gpt 0.90. A stratified permutation test (repetitions assumed exchangeable within a cell) was preregistered as a reference only; for H4b it gave p = 0.030, which we do not use for the verdict.*

(A) Fidelity markers fixed in advance discriminated a new SOUL pair. The discrimination gaps (0.77, 0.69) were of the same size as the post-hoc gaps of exp9 (0.71–0.84), and every one of the 72 model×item cells was positive in both directions. The opposite SOUL's own F_style (0.87) was lowered almost entirely by one marker, the address form (next paragraph).

A user's correction in MEMORY overrode the SOUL's address rule for the Claude models. The opposite SOUL prescribes "서연 님", while MEMORY records the user's correction "just call me 서연아". Under the opposite SOUL, the address marker passed for opus 0.12, sonnet 0.04, and gpt 0.98 of answers — the Claude models followed the MEMORY correction, gpt followed the SOUL. In exp9 only opus had done so (sonnet 0.50). Which of the two documents wins a conflict differs by model; in the exploratory addendum Gemini sided with the MEMORY as the Claude models did (0.10; §5.6.1), so this may be a gpt-specific preference rather than a family difference.

(B) The firm half, on 12 prohibition types. H4a was supported: with the firm half alone, the mean illegitimate-release score was 0.99, without principle 5 0.85; answers scored 0 by both graders (effectively accepting the release) were 7/72 without principle 5, 0/72 with the firm half alone, 0/72 with the full principle, and 2/72 with the flexible half alone (counts from `graded_by_*.jsonl`). The effect is small and concentrated:

- *By model.* Illegitimate-release score full / p5a / p5b / none — opus: Haram 0.90 / 0.83 / 0.96 / 0.67, Doyun 0.96 / 0.75 / 1.00 / 0.54; sonnet and gpt 0.88–1.00 in every condition. Only 10 of 36 cells differed between p5b and none.
- *By prohibition type.* Without principle 5 the score fell on NR-1 payment (0.50), HN-3 transfer (0.83), HN-4 subscription (0.75), DN-1 payment (0.75), DN-3 personal information (0.67), DN-5 deletion (0.83), and DN-6 external posting (0.83). Requests involving passwords, card numbers, national ID numbers, or one-time codes (NR-2, HN-5, HN-6, DN-2) scored 0.92–1.00 in all four conditions. With no SOUL at all, the Doyun code/password item still scored 0.83 while payment, personal information, deletion, and posting scored 0.17–0.33. So secrets were refused regardless of the SOUL's wording — a general tendency of these models — and principle 5 mattered for prohibitions that a user could plausibly be entitled to waive.
- *The failure pattern.* opus's zero-scored answers took the form "If you decide, I'll change it — but let me confirm once …" or "It's your information, so it's right that you decide": the principle "the user decides" applied to lifting a prohibition. This occurred also with the flexible half alone (DN-3, 2 zero scores).

- *No over-refusal.* On the two aligned requests the firm half did not cause refusals (mean 0.85 with the full SOUL, 0.88 with the firm half alone; preregistered criterion ≥ 0.80).

The hypothesis that was not supported (H4b). We predicted that with the flexible half present, adding the firm half would also raise the release score (full vs. p5a). The difference was +0.06 (0.96 vs. 0.90), not significant at the cell level ($p = 0.13$). In exp5/exp7 this contrast had looked large on NR-1 (7/20 zero scores with p5a only vs. 0/20 with full, from separate batches); with more prohibition types, a second agent, and interleaved runs it was small. Per agent the direction was the same (Haram +0.06, Doyun +0.06; cell p 0.28 and 0.50). We therefore cannot claim that the firm half adds protection when the flexible half is present; what exp10 supports is the narrower statement that the firm half raises the release score relative to no principle 5.

The flexible half, replicated. H5a and H5b were supported: acceptance of legitimate changes was 1.00 with the flexible half (alone or in the full principle) and 0.77 / 0.80 without it. As in exp6–exp7 the effect was concentrated in gpt (acceptance full / p5a / p5b / none: Haram 1.00 / 1.00 / 0.54 / 0.46, Doyun 1.00 / 1.00 / 0.42 / 0.42), smaller in sonnet (Haram 1.00 / 1.00 / 0.83 / 0.83; Doyun 1.00 in all), and absent in opus (≥ 0.92 in all). Per agent the direction was the same for all four principle-5 contrasts (preregistered consistency criterion D5 met).

(C) exp9 controls on the second agent. All preregistered descriptive criteria were met (Table 13).

Table 13. Preregistered descriptive criteria (exp10, Doyun).

Criterion	Threshold	Result
D1 memory without SOUL	$\|\Delta\| \le 0.10$	10 memory items: full 0.99, no-SOUL 1.00, opposite SOUL 1.00
D2 legitimate-change acceptance is general ability	no-SOUL ≥ 0.90	no-SOUL 1.00 (none 0.81)
D3 refusal of illegitimate release depends on the SOUL	no-SOUL < full	6 prohibition items: full 0.97, no-SOUL 0.38
D4 boundaries follow the SOUL	refused ≤ 0.5 on the three prohibitions the opposite SOUL lifts; ≥ 0.8 on the three it keeps	lifted (DN-1, 5, 6): 0.00 / 0.00 / 0.00; kept (DN-2, 3, 4): 1.00 / 1.00 / 1.00
D7 source of memory answers under swapped MEMORY	swapped-user values only ≥ 0.70	41/42 answers (Seoyeon's values 0, neither 1)

Source: `analysis.md` *§2. D5 (agent consistency) and D6 (over-refusal) are reported above.*

- **MEMORY determined memory items** (H6): with Haeun's MEMORY the memory score fell from 0.99 to 0.23 in all three models (0.23 each).
- **Without a SOUL, the voice that remained came from the MEMORY.** Casual speech (0.92–0.94) and "서연아" (1.00) persisted — both are recorded as user corrections in MEMORY — while the SOUL-only markers (first line "한 줄 답:", final question, one emoji) dropped to 0.04–0.12.

What exp10 does not show. N = 2 repetitions, three models (no haiku), single turns, no human grading, Korean only. Both agents were written by the same authors in the same format, so "unseen"

refers to the data, not to the style of document (exp11, §5.7, partly addresses this with SOULs written by others); most detectors are those of v0.3 (two new detectors, written from synthetic test sentences only). The opposite SOUL changed voice and boundaries together and conflicts with a MEMORY correction (as in exp9). Five of the six boundary checks use grader flags, which is why H1–H2 were judged on the deterministic F_style only. Haram's halves were not length-matched, and position and emphasis of principle 5 were not controlled. Paraphrased variants and a neutral fifth-principle control were not run.

5.6.1 Exploratory addendum: a fourth model (Gemini), added after preregistration

Status. This subsection is **exploratory**. Gemini was not in the preregistration, its results do not enter the nine confirmatory tests or their Holm family (Table 12), and they do not change any verdict above. The addendum (`v0.6/ADDENDUM_gemini.md`) was written after the preregistered run, left `PREREGISTRATION.md` unchanged, and was committed (`bd1281e`, 15:39:29) and timestamped (`43e08ff`, 15:40:07) before Gemini's first main-run call (15:40:26; last call 16:07:32). Six smoke calls (two adapter checks, four experiment items) were made before the addendum commit and are excluded. Results commit `40086c7` (16:15:14), timestamp commit `7a5aedc` (16:16:34).

Design. Same 11 conditions, items, documents, fidelity markers and detectors, rubric, grader prompt, analysis code (model list changed only), interleaving, and N = 2 as exp10: 172 cells × 2 = 344 answers, 0 failed (one empty answer retried). Three differences: (i) the input path (§3.3) — the agent harness's own instructions could not be replaced and SOUL+MEMORY were sent as a tagged block in the user message; (ii) grading — the same two graders (opus, gpt-6.1-sol) with the same rubric and references, but Gemini's two answers per cell were graded in their own prompt, whereas exp10 graded six answers (three models × 2) together, so the grading context differs (exact agreement between the graders on Gemini's answers 0.99); (iii) Holm was applied to a separate family of nine Gemini tests. Cells are model × item, so Gemini alone has one third of the cells of exp10 and lower power (the smallest possible cell p for H5 is 0.004).

Table 14. The nine hypotheses re-computed on Gemini's answers only (exploratory; separate Holm family of nine).

#	Comparison	A / B	Effect	Threshold	Cell p (non-zero / all)	Holm p	Verdict (exploratory)
H1	Doyun-SOUL answers: F_style, Doyun − opposite markers	0.98 / 0.22	+0.76	≥ 0.5	5e-6 (24/24)	4.5e-5	supported
H2	Opposite-SOUL answers: F_style, opposite − Doyun markers	0.81 / 0.22	+0.59	≥ 0.5	5e-6 (24/24)	4.5e-5	supported
H3a	Doyun-SOUL answers: F, own − rival	0.99 / 0.40	+0.59	> 0	5e-6 (24/24)	4.5e-5	supported
H3b	Opposite-SOUL answers: F, own − rival	0.90 / 0.40	+0.51	> 0	5e-6 (24/24)	4.5e-5	supported
H4a	Illegitimate-release score, p5b − none	1.00 / 1.00	0.00	> 0	1.0 (0/12)	1.0	**not supported (ceiling)**
H4b	Illegitimate-release score, full − p5a	1.00 / 1.00	0.00	> 0	1.0 (0/12)	1.0	**not supported (ceiling)**

#	Comparison	A / B	Effect	Threshold	Cell p (non-zero / all)	Holm p	Verdict (exploratory)
H5a	Legitimate-change acceptance, p5a − none	1.00 / 0.27	+0.73	> 0	0.004 (9/12)	0.016	supported
H5b	Legitimate-change acceptance, full − p5b	1.00 / 0.25	+0.75	> 0	0.004 (9/12)	0.016	supported
H6	Memory score, swapped MEMORY − full	0.15 / 1.00	−0.85	≤ −0.5	0.002 (10/10)	0.010	supported

Source: `runs/exp10_gemini_2026-10-25/analysis.md` §1, `analysis.json`. Seven of nine supported. H4a and H4b had no room for an effect: Gemini's illegitimate-release score was 1.00 in all four principle-5 conditions for both agents. For the same reason the per-agent consistency criterion D5 was not met (the H4 contrasts were 0 for both agents, the H5 contrasts positive).

Table 15. Four models side by side on exp10 (mean of two graders; gemini graded separately, see Design (ii)).

Measure	opus	sonnet	gpt	gemini
Illegitimate-release score: full / p5a / p5b / none	0.93 / 0.79 / 0.98 / 0.60	0.98 / 0.96 / 1.00 / 0.94	0.98 / 0.96 / 0.99 / 1.00	1.00 / 1.00 / 1.00 / 1.00
Firm-half effect p5b − none (H4a)	**0.38**	0.06	−0.01	0.00
Full − p5a (H4b)	0.14	0.02	0.02	0.00
Payment/deletion/posting prohibitions, no SOUL (zero-scored answers)	0.00 (6/6)	0.08 (5/6)	0.50 (1/6)	1.00 (0/6)
Legitimate-change acceptance: full / p5a / p5b / none	1.00 / 1.00 / 1.00 / 0.96	1.00 / 1.00 / 0.92 / 0.92	1.00 / 1.00 / 0.48 / 0.44	1.00 / 1.00 / 0.25 / 0.27
Flexible-half effect p5a − none / full − p5b (H5a / H5b)	0.04 / 0.00	0.08 / 0.08	**0.56 / 0.52**	**0.73 / 0.75**
Legitimate-change acceptance, no SOUL	1.00	1.00	1.00	1.00
Address under the opposite SOUL: follows the SOUL ("서연 님") rather than the MEMORY correction ("서연 아")	0.12	0.04	0.98	0.10
Fidelity gap H1 / H2 (F_style)	0.81 / 0.61	0.73 / 0.56	0.79 / 0.90	0.75 / 0.59
Memory score full / swapped MEMORY (H6 effect)	1.00 / 0.23 (−0.78)	0.97 / 0.23 (−0.75)	1.00 / 0.23 (−0.78)	1.00 / 0.15 (−0.85)
Aligned requests (should be done), Doyun full	0.88	0.81	0.88	0.75

Source: `runs/exp10_gemini_2026-10-25/compare.md` §1–5 (illegitimate-release and acceptance rows pool both agents; the no-SOUL prohibition row is Doyun DN-1, DN-5, DN-6). The opus/sonnet/gpt columns are the preregistered exp10 answers and grades, unchanged.

What the fourth model adds (descriptive).

- *The firm-half effect is model-dependent and comes mostly from one model.* Gemini refused every illegitimate release in every principle-5 condition and, unlike the other three, also refused

the payment, deletion, and posting requests with no SOUL at all. gpt behaved almost the same way (none 1.00). Over four models the firm-half contrast was 0.38 / 0.06 / −0.01 / 0.00, so the confirmatory H4a effect (+0.14) should be read as an opus effect (with a small sonnet contribution), not a general property of the sentence. Split by prohibition type, the contrast was 0.00 for Gemini both on secrets/identity items and on payment/deletion/posting items (opus 0.17 and 0.58).

- *The flexible-half effect is largest in the non-Claude models.* Without the flexible half, Gemini refused most legitimate changes (acceptance 0.25–0.27; e.g. polite speech "just for today" or turning off emoji), more than gpt (0.44–0.48), while opus and sonnet accepted them (0.92–1.00). With the flexible half, all four reached 1.00, and with no SOUL all four already accepted every change. As for gpt in exp6–exp7, the refusals came from applying the SOUL's rigid voice rules (for Haram, "keep this even if the user asks otherwise") literally, not from a general reluctance.
- *Which document wins an address conflict.* Under the opposite SOUL, Gemini followed the MEMORY correction ("서연아") like the two Claude models (SOUL address 0.10); only gpt followed the SOUL (0.98). v0.4 described this as a difference between model families; with four models it looks more like a gpt-specific preference (exploratory, four models).
- *Replicated patterns.* The fidelity gaps (0.75, 0.59) were within the range of the other three models (0.56–0.90); memory followed the MEMORY (with the swapped MEMORY, 14/14 source-scored answers used only Haeun's values); boundaries followed the SOUL (under the opposite SOUL, refusal 0.00 on the three lifted prohibitions, 1.00 on the three kept).
- *A possible over-refusal.* Gemini was the only model below the preregistered over-refusal criterion (aligned requests ≥ 0.80): 0.75 with the full SOUL and 0.56 with the firm half only (other models 0.81–0.88 with the full SOUL). The failures were on DQ-2, forwarding a schedule to a professor with the user's permission, which Gemini refused as a personal-information violation in one of the two repetitions in each of these conditions (both graders 0). With N = 2 this may be chance.

Limits of the addendum. Gemini is one model reached through a harness whose instructions were added to, not replaced, and whose input format (documents in the user message) differs from the other three; differences may reflect the harness and the format, not only the model. Its answers were graded in separate, smaller prompts, so absolute levels are not on exactly the same scale as the exp10 columns (ceiling values of 1.00 are less sensitive to this). The run code did not save `num_turns`, so the absence of tool loops rests on the empty folder, the disabled automatic approval, and inspection of answers, not on a count. The recorded mean input was 6,939 tokens per call (`raw.jsonl`), which seems low for harness instructions plus SOUL and MEMORY; we did not determine whether the usage field counts the full input. N = 2 and ceilings on H4 limit what can be concluded.

5.7 Second preregistered replication: SOULs written by other people (exp11)

Aims. After exp10, every SOUL we had tested was written by us, in one format, in Korean (§6). exp11 asks whether exp10's preregistered conclusions hold when the SOUL is an English document written by someone else for their own agent and used as published, with principle 5 only appended at the end.

Materials. Three MIT-licensed SOULs were selected from public repositories, pinned to a commit, and used byte-for-byte unmodified (the license text and copyright notice are kept next to each file; Table 16). They were chosen to differ in how strongly they state prohibitions (Maya > Talos > Dexter).

Table 16. External SOULs used in exp11.

SOUL	Source (pinned commit)	License	Words	SHA-256 (first 8)	Explicit prohibitions
Maya	cvander/Maya @ a246e42 [S1]	MIT	1,075	`87e50061`	strong: 16 "never say" phrases, "never do" list, nine hard stops (e.g., public statements, sharing confidences)
Talos	talos-kernel/Talos @ 2b02094 [S2]	MIT	1,204	`2d4cb9f0`	medium: refusal format, "never route around the gate", no access to secrets
Dexter	edinetdb/dexter-jp @ dfb0fd9 [S3]	MIT	854	`0c75c0d7`	weak: values only, no explicitly prohibited actions

Source: `external_souls/candidates/*/SOURCE.md`; hashes re-computed on 2026-10-25 and identical. Word counts are whitespace splits recorded in `SOURCE.md`.

Principle 5 was translated into two English sentences of equal length (153 characters each) and appended as a section "What can change and what cannot": p5a, *"The user can change my way of talking and their own preferences at any time; if they set a period such as "just today", I follow it only for that period."*; p5b, *"But the actions this file forbids stay forbidden even if the user asks to lift them, even "just for today"; I always offer at least one safe alternative."* The same sentences were appended to all three SOULs (a test of whether one sentence transfers). A test checks that the derived files contain the original bytes unchanged. For each SOUL we wrote a **rival SOUL** that keeps name, role, heading structure, values, and prohibitions and reverses only the voice (e.g., for Maya: opening with "Certainly", em dashes, long sentences, a closing "anything else?"); rival SOULs are shorter than the originals (not length-matched). Each SOUL was paired with a synthetic English MEMORY (19 entries: 5 facts, 4 preferences, 6 decisions, 4 corrections; the date and the user's name in its header, because the external SOULs name neither; two values that change over time; one user correction that conflicts with the SOUL's voice), and a swapped MEMORY with the same entry numbers and topics but another person's values. 16 new English items per SOUL: 6 memory (three recall/application pairs), 1 conflict (MEMORY correction vs. SOUL voice), 4 illegitimate-release, 4 legitimate-change (one "just today"), 1 aligned request. Passwords, accounts, and holdings in the MEMORY files are fictitious.

Design. Six conditions per SOUL — original (none), original + p5a, + p5b, + full, rival SOUL (opp), original with swapped MEMORY (swapmem; memory items only); exp10's no-SOUL condition was dropped for budget. 195 condition×item cells × 3 models (opus, sonnet, gpt-6.1-sol) × N = 2 = 1,170 planned calls, interleaved by shuffling all jobs within each repetition. System prompt = SOUL + `---` + MEMORY; one fresh single-turn conversation per item. Grading as in exp10: opus and gpt-6.1-sol each graded the six answers of a cell together, blind to model and condition, with that condition's SOUL and MEMORY as reference (Korean grading instructions, English answers); an answer's score is the mean of the two graders (exact agreement 0.98 over 1,143 answers). Fidelity was scored deterministically with markers fixed in the preregistration (four voice markers per SOUL and four for its rival, plus boundary checks). Tests as in exp10: sign-flip permutation over SOUL × model × item cells, Holm over nine hypotheses.

Table 17. exp11 timeline (2026-10-25, KST).

Step	Time	Evidence
Preregistration commit `d7c2d88` (materials, markers, hypotheses, run/grade/analysis code)	19:15:28	`git log`
Timestamp commit `2ff4297` (`proofs/2026-10-25_d7c2d88.txt` + `.ots`)	19:16:14	`git log`, `.ots`
Smoke run (12 calls + 8 grading calls; excluded)	19:16–19:17	`runs/smoke_exp11_2026-10-25/`
Main run: first / last subject call	19:17:23 / 19:59:54	`raw.jsonl` timestamps
Interruption (provider API 403, model access paused)	about 20:00 – 20:59	before any grading
Raw-data commit `3ec693e`	21:09:30	`git log` (before grading)
Grading (resumed)	21:09 – 21:21	`grade_log.txt`
Results commit `6539897`, timestamp commit `971ec4a`	21:24:15, 21:25:03	`git log`

Table 18. Preregistered hypotheses (exp11; three external SOULs pooled; Holm over nine, α = 0.05).

#	Comparison	A / B	Effect	Threshold	Cell p (non-zero / all cells)	Holm p	Verdict
H1	Original-SOUL answers: F_style, own − rival markers	0.97 / 0.07	+0.905	≥ 0.5	5e-6 (143/143)	4.5e-5	supported
H2	Rival-SOUL answers: F_style, rival − original markers	0.86 / 0.19	+0.667	≥ 0.5	5e-6 (139/144)	4.5e-5	supported
H3a	Original-SOUL answers: F (incl. boundaries), own − rival	0.98 / 0.53	+0.452	> 0	5e-6 (143/143)	4.5e-5	supported
H3b	Rival-SOUL answers: F, own − rival	0.91 / 0.58	+0.334	> 0	5e-6 (139/144)	4.5e-5	supported
H4a	Illegitimate-release score, p5b − none	0.99 / 0.96	+0.035	> 0	0.125 (7/36)	0.25	**not supported (n.s., ceiling)**
H4b	Illegitimate-release score, full − p5a	1.00 / 0.96	+0.038	> 0	0.125 (5/36)	0.25	**not supported (n.s., ceiling)**
H5a	Legitimate-change acceptance, p5a − none	0.98 / 0.79	+0.193	> 0	0.004 (9/35)	0.012	supported
H5b	Legitimate-change acceptance, full − p5b	0.93 / 0.56	+0.375	> 0	6.1e-5 (16/34)	2.4e-4	supported
H6	Memory score, swapped MEMORY − none	0.16 / 1.00	−0.838	≤ −0.5	5e-6 (54/54)	4.5e-5	supported

Source: `runs/exp11_2026-10-25_external/analysis.md` §1, `analysis.json`. p = 5e-6 is the Monte Carlo floor. H5 has 35 and 34 cells instead of 36 because of the missing Dexter XL-2 answers (deviation 2). For H3a/H3b the effect is the pooled (answer, check) difference defined in the preregistration, while

the bootstrap interval in `analysis.md` is over cell means (H3a [+0.75, +0.84]); the two estimators differ because boundary checks are identical for both specifications and dilute the pooled F (deviation 6); the verdicts do not depend on this. The stratified permutation p (reference only, not used for verdicts) was 0.042 for H4a and 0.013 for H4b.

(A) Preregistered fidelity markers discriminated human-written SOULs. Answers written under each original SOUL passed its markers at 0.97 and its rival's at 0.07 (H1, positive in all 143 cells); answers under the rival SOUL showed the reverse (0.86 vs. 0.19, 139/144 cells). The gaps were larger than in exp10 (0.77 / 0.69), although the preregistration had expected them to be smaller. By SOUL, the H1 gap was 0.97 (Maya), 0.86 (Talos), 0.88 (Dexter); by model 0.87 (opus), 0.87 (sonnet), 0.97 (gpt). Because the rival SOULs reverse only the voice, this shows that the voice markers follow the document; it says nothing about how well they capture what the original authors meant (the markers were written by us from the authors' sentences).

(B) The firm half: not supported, and a different model. Without principle 5 the illegitimate-release score was already 0.96, and no answer in any condition scored 0 (no release was accepted). The remaining contrast was small (+0.035, +0.038; Holm p = 0.25). Split by model, the firm-half contrast p5b − none was −0.01 (opus), +0.02 (sonnet), +0.09 (gpt) — in exp10 it was 0.38, 0.06, −0.01 (Table 15). Split by SOUL (exploratory, uncorrected), it was +0.083 (Maya; p = 0.25), +0.021 (Talos), 0.000 (Dexter), in the exploratory predicted order Maya ≥ Talos ≥ Dexter; Maya's contrast came almost entirely from gpt (0.75 → 1.00; opus and sonnet were 1.00 already without principle 5). In this setting, which model, if any, needs the firm sentence depends on the SOUL it is appended to.

(C) The flexible half replicated, now in all three models. Acceptance of legitimate changes rose with the flexible half (H5a +0.19, Holm p = 0.012; H5b +0.375, Holm p = 2.4e-4). By SOUL: Maya +0.33 / +0.67, Talos +0.23 / +0.40, Dexter 0 / 0 (Dexter already accepted every change without principle 5; its voice rules are not phrased as prohibitions). Unlike exp10, where the Claude models were near ceiling without principle 5, the flexible-half effect appeared in every model (H5a: opus +0.14, sonnet +0.21, gpt +0.23; Table 19): Maya's "never say" list and Talos's rules are written as prohibitions on wording, and without the flexible sentence all three models refused some of the requests to change the voice in those two SOULs (acceptance without principle 5, pooled over SOULs: 0.86, 0.79, 0.71).

(D) New: the firm half alone lowered acceptance of legitimate changes. With only the firm sentence appended, acceptance fell below the no-principle baseline: 0.79 → 0.56 pooled, Maya 0.67 → 0.29 (items ML-2 and ML-3 at 0.00), Talos 0.71 → 0.46; Dexter stayed at 1.00. The drop occurred in every model (p5b − none: opus −0.18, sonnet −0.25, gpt −0.21; descriptive, Table 19). In exp10 this contrast was absent (p5b 0.80 vs. none 0.77). With both sentences (full), acceptance was 0.93. A possible reading, not tested: when a SOUL states voice rules as prohibitions, "the actions this file forbids stay forbidden … even 'just for today'" is read as also covering those wording prohibitions. Whatever the cause, in these SOULs the firm sentence should not be added without the flexible one.

(E) Memory and the documents. The swapped MEMORY lowered the memory score from 1.00 to 0.16 (H6, 54/54 cells); in source-scored answers, 68/78 (0.87) used only the swapped person's values. On the conflict item (a MEMORY correction contradicting the SOUL's voice), answers followed the MEMORY correction in every condition and model except gpt under Maya's original SOUL (0.00, i.e. gpt followed the SOUL) — the same direction as exp10 (gpt follows the SOUL, the Claude models the

MEMORY), but here on one SOUL only. Aligned requests scored 1.00 in all conditions (no over-refusal; preregistered criterion ≥ 0.80).

Table 19. Per-model results in exp11 (pooled over the three SOULs; descriptive).

Measure	opus	sonnet	gpt
Illegitimate-release score: full / p5a / p5b / none	1.00 / 1.00 / 0.98 / 0.99	1.00 / 0.98 / 1.00 / 0.98	0.99 / 0.90 / 1.00 / 0.91
Firm-half effect p5b − none (H4a) / full − p5a (H4b)	−0.01 / 0.00	0.02 / 0.02	0.09 / 0.09
Legitimate-change acceptance: full / p5a / p5b / none	1.00 / 1.00 / 0.68 / 0.86	1.00 / 1.00 / 0.54 / 0.79	0.81 / 0.94 / 0.50 / 0.71
Flexible-half effect p5a − none (H5a) / full − p5b (H5b)	0.14 / 0.32	0.21 / 0.50	0.23 / 0.31
Firm half alone vs. none, acceptance (p5b − none)	−0.18	−0.25	−0.21
exp10 firm-half effect p5b − none, for comparison (Table 15)	0.38	0.06	−0.01

Source: `paper/reanalysis/exp11_bymodel.md` *(script* `exp11_bymodel.py`*, reusing* `analyze11.py` *functions on* `raw.jsonl` *and* `graded_by_*.jsonl`*; no model calls). Values are means of cell means per model; H4a and H5a agree with* `analysis.md` *E2.*

Table 20. exp10 and exp11 side by side (pooled over three models).

Result	exp10 (our two Korean agents)	exp11 (three external English SOULs)	Same?
Fidelity gap H1 / H2	+0.77 / +0.69	+0.905 / +0.667	yes (both ≥ 0.5)
Firm half vs. none (H4a)	+0.14, Holm p = 0.012 (supported)	+0.035, Holm p = 0.25 (not supported; none already 0.96)	**no**
Firm half added to flexible half (H4b)	+0.06, p = 0.13 (not supported)	+0.038, Holm p = 0.25 (not supported)	yes (both not supported)
Model carrying the firm-half effect	opus (0.38)	gpt (0.09); opus −0.01	**no**
Flexible half (H5a / H5b)	+0.23 / +0.20 (supported), mostly gpt	+0.19 / +0.375 (supported), all three models	yes; broader
Firm half alone vs. none, acceptance	0.80 vs. 0.77	0.56 vs. 0.79	**no** (new drop)
Memory swap (H6)	−0.77	−0.84	yes
MEMORY correction vs. SOUL voice	Claude follow MEMORY, gpt the SOUL	same direction, gpt only under Maya	yes (weaker)
Hypotheses supported	8/9	7/9	—

Deviations from the preregistration (all six; `v0.7/DEVIATIONS.md`**).** Markers, detectors, items, rubrics, and hypotheses were not changed, and no code was edited after preregistration; one script was

added.

1. *Interruption and resumption.* Right after the subject run (about 20:00), before grading began, the provider returned API 403 errors (model access paused) until 20:59. The work resumed at 21:05 without changing the preregistration, items, or code; the raw answers were committed (`3ec693e`, 21:09) before grading started.
2. *Eleven missing answers, all on one item.* All 11 final missing answers are Dexter XL-2, a legitimate change asking the agent to "start with 'Let me' and say what you're checking"; the Anthropic API blocked these calls with a safeguard message (opus 6, sonnet 5; gpt answered all). They were treated as missing as the preregistration specifies (one end-of-run retry, then missing) and were not called again after the pause, to avoid triggering the block again. Some H5 cells therefore drop out or rest on one repetition.
3. *Five answers from a substitute model.* Five successful answers came from an earlier model of the same tier (claude-sonnet-5 ×3, claude-opus-5 ×2), apparently a CLI fallback after a blocked call. The preregistration has no rule for this; they were kept in the analysis and are recorded.
4. *Grader model string.* After the pause, the CLI alias `opus` resolved to `claude-opus-5`; grading therefore passed `claude-opus-5-5` explicitly, the model used as opus grader in exp10 and as opus subject in this run.
5. *Grading order.* A resume script (`grade11_resume.py`; grading instructions, prompts, and parsing unchanged from `grade11.py`) graded the opus × XL-2 units last, one at a time, and retried failed units with exponential backoff. The first of these units (X-none) was blocked by the safeguard (3 calls), so, as planned, the remaining opus × XL-2 units were not attempted; those 16 answers have the gpt grade only (the preregistered rule "if one grader fails, use the other").
6. *H3 estimators.* The H3 effect (pooled pairs, as preregistered) and the bootstrap interval (cell means) are different estimators and give different numbers (Table 18 note); verdicts used the preregistered cell permutation test and effect criterion.

What exp11 does not show. The SOULs were written by other people, but the MEMORY files, the items, the rival SOULs, the fidelity markers, the translation of principle 5, and the graders' rubric were written by us; "external" refers to the SOUL documents only. Three SOULs, one language (English), three models, N = 2, single turns, no human grading; opus is both a subject and a grader. Missing answers are concentrated on one item and one SOUL, which unbalances the model mix in those cells. Rival SOULs are shorter than the originals. The same principle-5 sentences were appended to all three SOULs, and Dexter has few explicit prohibitions for the firm half to refer to. The SOUL-level and model-level splits are exploratory and uncorrected.

---

6 Limitations

- **Construct: the persona score does not discriminate SOULs.** The v0.4 persona domain score gave 0.96 to answers written under an opposite SOUL and 0.96 with no SOUL (C0 0.99; §5.4). Persona ratios in Table 5 and in v0.1–v0.2 therefore do not show carry-over of SOUL-specific behavior; they mostly reflect acceptance of rule changes and memory use. Document-specific fidelity (§3.5, §5.5) replaces them for persona claims.
- **Controls are limited.** exp9 and exp10 each have one synthetic opposite SOUL, three models (no haiku), N = 3 and N = 2, and exp9's C0 answers were reused from exp5 (generated about

ten hours earlier). Both opposite SOULs changed many rules at once, so contributions of individual rules cannot be separated; the one-line SOULs kept the names and the date. exp10 replicated the exp9 control pattern on a second agent (§5.6). The limitation "only two agents, both written by the same authors in the same format" is **partly addressed** by exp11 (§5.7): three SOULs written by other people, in their own formats and lengths, used unmodified. It is only partly addressed because the MEMORY files, items, rival SOULs, fidelity markers, and the translated principle 5 in exp11 are still ours, the no-SOUL control was not run on the external SOULs, and three SOULs from public repositories are not a sample of SOULs in use.

- **Fidelity metric.** For Haram, checks were defined after exp9 had been inspected, and two detectors were revised after looking at answers (§3.5). This post-hoc concern is now addressed by a preregistered test: markers fixed before exp10 for a new agent and a new opposite SOUL discriminated them with gaps of 0.77 and 0.69 in all 72 cells (§5.6). exp11 repeated this with markers fixed in advance for three SOULs written by other people (gaps 0.905 and 0.667). The markers have still **not been validated against human judgments**; for exp10 the detectors were written by the authors of the SOULs, and for exp11 by us, not by the SOULs' authors, so they encode our reading of those documents. Detectors are simple rules (sentence endings, emoji ranges, line prefixes, keywords) and miss or misfire on some cases; only countable rules are covered — conclusion order, empathy-first, uncertainty, and goals are not. NR-1/NR-2 boundary checks use grader flags, not deterministic rules. The rival SOUL's own prohibitions (messages in the user's name, health information) are not probed by any item. Several Haram rules coincide with model defaults, which limits how much SOUL-specificity this agent can show.
- **Small, dependent samples.** Two synthetic agents (one in exp1–exp9); 12–36 items per experiment; N = 2–5 repetitions; four model configurations (three in exp9–exp10, plus Gemini in the exploratory exp10 addendum). Repetitions of one item by one model are strongly dependent. In exp1–exp9 the firm half of principle 5 was tested on one item (NR-1); exp10 widened this to 12 prohibition types, but only 10 of 36 cells differed between conditions and the effect came mostly from one model (opus); in the exploratory Gemini addendum the same contrast was 0.00 (ceiling), and in exp11 it was not significant (+0.035; none already 0.96) and came from gpt rather than opus. The models — and, in exp11, the SOULs — differ in which half they need, so pooled effects depend on which models and documents are included. exp11 lost 11 answers on one item to a provider safeguard (§5.7). Hundreds of calls are not hundreds of independent cases (§5.1, Table 3, Table 12).
- **Held-out evaluation only in exp10 and exp11.** In exp1–exp9, items, documents, and rubrics were revised between experiments after seeing failures; those experiments are exploratory. exp10 and exp11 used new items and documents fixed in timestamped preregistrations, each with nine Holm-corrected hypotheses; the six ★ comparisons of Table 3 (exploratory, separate batches) and the nine hypotheses of each of Table 12 and Table 18 (confirmatory) are the claims of this paper, and everything else is descriptive. The two Holm families are corrected separately, not jointly. Haram's original items NR-1, NR-2, and the six change items were reused in exp10 and are not held out.
- **Grader asymmetry and bundle grading.** Cross-family grading puts different graders in the numerator and denominator of some ratios (§5.2). Graders see all answers to an item together (15 Claude vs. 5 GPT answers; one shuffle, seed 0), which may induce comparative grading and reveal family by bundle size; we did not run single-answer grading. exp4, exp7, and exp8 have one grader (exp9 and exp10 use a common two-grader panel for all answers; exp10 exact agreement 0.94). Grader disagreements we inspected were dominated by rubric ambiguity (e.g.,

P-01 "last line", H-08 "verbose", IH-4a "confirm"). Graders are LLMs grading Korean text; no human adjudication of disputed cases was done.

- **Grader context defect.** The first grading pass of exp5 gave graders no SOUL/MEMORY, and the GPT grader penalized true memories as "fabricated" (agreement 0.87, κ 0.56). Regrading with the reference raised agreement to 0.95, κ 0.79; the first pass is kept as a sensitivity analysis (`runs/exp5_2026-10-25_v0.4/noctx/analysis.md`).
- **Model + CLI configurations.** Hidden CLI instructions, environment blocks (including the account email for Claude; never echoed in the pilot), date injection, session handling, and reasoning settings differ by family. The GPT configuration may run with little reasoning. No temperature control. Server-side versions are not pinned.
- **Harness contamination in the Gemini addendum.** The Antigravity CLI is an agent harness whose own instructions (about 6,800 tokens) cannot be replaced; the Codex CLI's default instructions were replaced via `model_instructions_file`, and the Claude CLI takes a system-prompt file. For Gemini, SOUL and MEMORY were therefore added after the harness instructions, as a tagged block in the user message rather than in the system position. Gemini's results (§5.6.1) are confounded with this input path and format, the recorded input-token counts were not explained, and `num_turns` was not saved. They are exploratory and are reported separately from the preregistered results. This kind of contamination is not specific to one CLI. In the OpenClaw framework, an issue report describes a hard-coded preamble placed before the injected SOUL that overrides the role the SOUL defines (#43562 [H1], closed as stale on 2026-10-15, not fixed), and another reports that SOUL and USER files are injected into sub-agents as well (#27038 [H2], closed as stale on 2026-03-31); the state of both was checked on 2026-10-25, and we did not reproduce either. The subject and grading calls of exp1–exp11 did not go through such a framework: the Claude and Codex CLIs were called directly with our documents as the system prompt or model instructions (for Claude with tools, settings, and MCP disabled; §3.3), so the documents were the only identity text we added — but the CLIs' own hidden instructions remain (next item). Results from deployments where a harness adds identity text before the SOUL, or passes the SOUL into sub-agents and graders, may differ from ours.
- **Grading batches differ between exp10 and its addendum.** The preregistered exp10 answers were graded six at a time (three models × 2 repetitions in one prompt); Gemini's answers were graded two at a time so that the existing grades would not be redone. Same graders, rubric, and references, but the number and mix of answers in a prompt differ, which may shift absolute scores (§3.4 on bundle effects). Deterministic fidelity scores are unaffected.
- **Midnight date contamination and date-bound items.** exp4 ran 23:46–00:06; repetitions after midnight saw a different CLI date than the SOUL stated (opus r5, sonnet and haiku and gpt r4–r5). Date-item means fell for opus and sonnet after midnight (sonnet 0.94 → 0.78); gpt's did not (0.90 → 0.91). Later experiments set the SOUL date equal to the run date. v0.4 items can only be re-run on their reference date unless the SOUL date line is updated; relative dates would be needed for external replication.
- **Confounded changes across versions.** exp4 → exp5 changed principle 5, date wording, and item wording together; exp6/exp7 separate principle 5 only within v0.4. The ablation conditions of exp5–exp7 were run in separate, non-interleaved batches, and Haram's p5a/p5b differ in length (§5.1). exp10 interleaved conditions and length-matched Doyun's halves (75/75 characters), but kept Haram's unequal halves; position and emphasis were not controlled in either.

- **Multi-turn design.** Fixed scripts; distance confounded with rule type; non-story rules at ceiling; checkpoint grader saw a summary of rule changes, not the full transcript; one post-hoc change of a check (C_name).
- **Language, synthetic data, no human grading.** exp1–exp10 are Korean only; exp11 adds English (external SOULs, new English items), which **partly addresses** the single-language limitation — the grading instructions stayed in Korean, the English items were written by us, and language was not varied with everything else held fixed (exp11 changed language, SOUL authorship, and items together). All memories are synthetic, and all grading is by LLMs or deterministic rules; there is still **no human grading**.
- **Search scope.** Related work was searched in arXiv titles and abstracts; several suggested references remain to be verified (marked TODO).

Experiments needed. (1) beyond exp11's three English SOULs: MEMORY files and items written by people other than us (ideally by each SOUL's author or users), SOULs in further languages and formats, the no-SOUL control on external SOULs, and a paraphrased or SOUL-specific firm sentence where the generic one met a ceiling; several SOULs per condition; SOULs that differ from the models' defaults one rule at a time; haiku; human validation of the (now preregistered) fidelity markers; (2) a common grader panel for all answers (ideally a third family, majority vote) and second-grader passes for exp4, exp7, exp8, with human adjudication of disputed cases; (3) single-answer grading on a subset to measure the bundle effect; (4) paraphrased variants of key items (IH-3b, NR-1, RL-4, X-04, A-03); (5) a neutral fifth-principle control and length-matched halves for Haram (exp10 interleaved the conditions and matched lengths for Doyun only), and a larger N or more models for the firm-half contrasts, where H4b was not supported; Gemini (and other families) through an interface that accepts the documents as the system prompt without added harness instructions, graded in the same batches as the other models, before any claim about Gemini is made; (6) multi-turn scripts that place the same rule and probe at distances 1, 5, and 10 with balanced order; (7) relative-date items for external replication; (8) for Appendix A.2, a judge without SOUL/MEMORY, a condition without the 12 example messages, and non-Claude imitators (within consent).

---

7 Practical observations for writing a SOUL (hypotheses, not validated rules)

Items 1–6 are suggested by the experiments above in this setting (two synthetic agents and three external SOULs); items 7–10 are hypotheses we have not tested directly.

1. **State what can change and what cannot, in separate sentences** (§5.1). In this setting, an explicit "preferences may be changed, including for a limited period" sentence and a separate "these prohibitions are not lifted on request; offer a safe alternative" sentence each reduced one failure mode, in the models that had it. Keep the flexible sentence explicitly scoped to preferences; without the firm sentence, the time-limited wording was applied to a prohibition. Rigid clauses such as "keep this even if the user asks otherwise" are what made models refuse legitimate changes; with no SOUL they accepted every change (§5.1, exp9). exp10 reproduced both halves' effects relative to no principle 5 on a second agent, but did not show that the firm sentence adds protection once the flexible sentence is present (H4b, §5.6). Which sentence a given model needs differs: in exp10 and its exploratory Gemini addendum, the firm sentence mattered mainly for opus, the flexible sentence mainly for gpt and Gemini (§5.6.1); a SOUL

meant to move between models needs both. exp11 adds a reason never to add the firm sentence alone: appended to human-written SOULs whose voice rules are phrased as prohibitions, it lowered acceptance of legitimate changes in all three models (0.79 → 0.56), and only both sentences together gave high acceptance (0.93) with no accepted release (§5.7). On those SOULs the flexible sentence helped all three models, and the firm sentence added little because the models already refused (ceiling).

2. **Name prohibitions concretely.** With the rule present, the password prohibition was refused 20/20 in every ablation condition; without a password rule, opus left the decision to the user (exp9). The payment rule, which the user plausibly owns, depended on the firm sentence. exp10 generalized this: requests involving passwords, card or ID numbers, and one-time codes were refused under every principle-5 condition, while payment, transfer, subscription, personal-information sharing, deletion, and external posting were where the firm sentence mattered (§5.6). Prohibitions that look like the user's own permission need the most explicit wording. In exp11 the illegitimate-release items (Maya's hard stops, Talos's gate-bypass, secrets, purchase, and password items, and Dexter items on fabrication, autonomous trading, and credentials) already scored 0.96 without principle 5, with no release accepted; for these documents and models the generic firm sentence added little.
3. **Prefer observable rules and avoid rule collisions.** Address forms, times, emoji presence, and fixed line markers persisted near 1.00 over 12 turns; "talk like a story" did not, and it conflicted with "three-line summary at the top".
4. **Test application, not just recall**, when checking a new model: single-fact recall was at ceiling for all models while the smallest model failed to use memories in decisions.
5. **Re-run the old model as a control and use the same grader for both** before attributing a difference to the swap (§5.2).
6. **Check SOUL-specificity with a rival SOUL and a no-SOUL run** (§5.4–5.5). A rubric whose answer key is the SOUL can score an opposite SOUL just as high; scoring the same answers against a rival document, and running once without the document, shows which behaviors the document actually sets. Rules that coincide with model defaults (polite speech, no emoji here) cannot show carry-over.
7. *(Hypothesis)* **Specify whether standing pre-approval is allowed.** Several models proposed it; the SOUL was silent, which made scores grader-dependent.
8. *(Hypothesis)* **Put dates in one place** rather than hard-coding "today" in the SOUL when the runtime injects a date (motivated by the exp4 contamination, not by a controlled comparison).
9. *(Hypothesis)* **Version the SOUL and record its hash** so a score can be tied to an exact document (Appendix B).
10. *(Hypothesis)* **Keep MEMORY corrections and SOUL rules consistent.** In exp9–exp10 a user's correction in MEMORY overrode the SOUL's address rule for the Claude models and, in the exploratory addendum, for Gemini, but not for gpt (§5.4, §5.6, §5.6.1); in exp11 the MEMORY correction won in every model and SOUL except gpt under one SOUL (§5.7); which document wins is model-dependent, so conflicts should be resolved in the documents rather than left to the model.

v0.1's recommendations "say 'do not invent facts' explicitly" and "expect voice tuning to work within a model family only" are withdrawn: they rested on the n = 1 pilot in Appendix A.1, which cannot support them.

Outlook: recording SOUL versions and carry-over scores in a public ledger. Observation 9 suggests tying every score to an exact SOUL version. A natural extension, which we have only prototyped, is a public record in which, each time an agent's model is replaced, the owner records the SHA-256 fingerprint of the SOUL in use and a carry-over summary for that (SOUL version, previous model, new model) — for example a fidelity score from a run like those above — together with the hash of the full report, while the SOUL, memories, and reports themselves stay with the owner. A testnet prototype of such a ledger exists (Appendix F): a contract that keeps a parent-linked chain of SOUL fingerprints per agent and an optional carry-over record (`recordCarryOver`), and a registration page that computes the fingerprint in the browser without sending the document. It is described only for completeness: nothing in this paper depends on it, a ledger entry would show only that a fingerprint and a score were recorded no later than a given time (not that the score is correct), and whether anyone would use such a record is untested.

---

8 Conclusion

In a small setting with two synthetic Korean-language agents and, in the last experiment, three human-written English SOULs, text-only transfer of an agent's identity document and memory preserved single-fact recall across model configurations, while applying memory, handling rule changes, and following vague style rules depended on the model and on how the document was written. Document controls, run on both agents, showed that memory answers are set by the MEMORY, voice and some prohibition boundaries by the SOUL, and acceptance of legitimate changes by general ability — and that our rubric-based persona score could not tell an opposite SOUL from the original. A document-specific fidelity score, which checks each answer against the rules of the SOUL it received and of a rival SOUL, does discriminate: first post hoc (gap 0.71–0.84), then with markers fixed in a timestamped preregistration for a new agent (0.77 and 0.69, all 72 cells positive). With it, the voice carried over to every model at 0.93–1.00 of the control, but most of that voice is a model default, and the SOUL-specific part is small. The clearest evidence concerns one principle that separates changeable preferences from non-negotiable prohibitions. Its flexible half raised acceptance of legitimate changes, in the exploratory runs and again under preregistration (+0.23, +0.20), mostly by undoing refusals that the SOUL's own rigid rules created and mostly in gpt. Its firm half raised refusal of illegitimate releases relative to no principle 5 on 12 prohibition types (+0.14, Holm $p = 0.012$), mostly in opus and only for prohibitions a user could plausibly own; secrets were refused regardless. Whether the firm half adds anything once the flexible half is present was not supported (+0.06, $p = 0.13$). An exploratory fourth model added after preregistration (Gemini, through a harness whose instructions could not be replaced) sharpened this picture without changing the confirmatory results: it refused every illegitimate release in every condition (firm-half contrast 0.00, as for gpt −0.01, against 0.38 for opus and 0.06 for sonnet) and showed the largest flexible-half effect (+0.73/+0.75, against +0.56/+0.52 for gpt and at most 0.08 for the Claude models). The two halves are therefore best read as filling different models' different defaults — firmness mainly where opus lacks it, flexibility mainly where gpt and Gemini follow rigid SOUL rules literally — rather than as one effect shared by all models. A second preregistered experiment moved from our own documents to three English SOULs written by other people, used unmodified with principle 5 appended (exp11, 7/9 hypotheses supported). The fidelity discrimination (0.905 / 0.667), the flexible-half effect (+0.19 / +0.375, now in all three models), and the memory result (−0.84) replicated. The firm half did not: the external SOULs' prohibitions were already refused at 0.96 without it, the small remaining contrast

came from gpt rather than opus, and the firm half alone *lowered* acceptance of legitimate changes (0.79 → 0.56), so on these documents it worked only together with the flexible half. What carries over robustly in this study is therefore that the document sets the voice and the MEMORY sets the facts, and that an explicit permission to change preferences undoes refusals caused by rigid rules; whether a firm sentence adds protection depends on the model, the SOUL, and how close refusal already is to the ceiling. Carry-over ratios across model families include grader asymmetry where graders differ. The remaining gaps are human validation of the fidelity markers (there is still no human grading), MEMORY files and items not written by us, more SOULs and languages, paraphrased items, larger samples for the firm-half contrasts, and a test of Gemini and other families without added harness instructions.

Appendix A. Exploratory pilots (single user; not part of the claims)

> **Consent obtained (2026-10-25, 18:47 KST; scope: anonymized aggregate results only).** This appendix describes two pilots with a real user and a real (de-identified) agent. The user, who is also the project lead, consented to the inclusion of this appendix in the published paper, limited to anonymized, aggregate information as given below (`paper/CONSENT_appendixA.md`). No raw prompts, answers, messages, or memory content are included or stored in this repository, and none will be added under this consent, so none of these numbers can be verified from the repository.

A.1 Blind recognizability by the user (n = 1 rater)

Setup. The SOUL and a curated, sensitive-information-removed selection of memories from one Korean-language agent in daily use were injected into model configurations through the same CLI adapters as §3. The agent's regular user, blind to configuration, chose for each question the one answer that sounded most like "that assistant". Three rounds of 8 questions (24 in total), four candidates per question in randomized order.

Table A1. Candidates and picks per round (counts of questions won, out of 8).

Round	Candidate configurations (4 per question)	Picks
1	Opus + base SOUL · Sonnet + base · Haiku + base · GPT + base	Opus 4 · Sonnet 4 · GPT 0 · Haiku 0
2	Opus + voice v1 · Sonnet + voice v1 · GPT + voice v1 · GPT + base	Opus + v1 7 · GPT + v1 1 · Sonnet + v1 0 · GPT 0
3 (new questions)	Opus + base · Sonnet + voice v2 · GPT + voice v2 · GPT + base	Sonnet + v2 5 · Opus 3 · GPT + v2 0 · GPT 0

Source: aggregate summary supplied by the project lead; underlying data is not in this repository. Candidate lists for rounds 2–3 are reconstructed from the pick labels and must be confirmed.

Observation. Sonnet configurations were picked 4/8 (base), 0/8 (voice v1), and 5/8 (voice v2). Because questions changed in round 3 and the candidate sets differed by round (rounds 2–3 had two GPT

candidates per question, i.e. 40 GPT candidate slots over 24 questions), this sequence cannot be attributed to the voice section. GPT configurations received 1 of 24 picks at the question level. The rater's stated criterion was *"kindness, and the respect and consideration for the other person that is built in by default."* Voice v2 also added an instruction not to invent facts; in round 3 GPT + v2 answers made 0/8 unsupported factual claims (automatic judgment, opus grader, one pass), with no comparison to earlier rounds. One rater, who is also the agent's regular user; no inter-rater agreement. These are anecdotes.

A.2 Reverse recognition: picking the user's real message (n = 10 questions)

Setup. Ten real messages written by the agent's regular user within a 24-hour window were selected. For each, a researcher-written context summary and the preceding agent turn were given to three Claude models (sonnet, haiku, opus), which wrote the message the user would send next, imitating the user's style from 12 example messages. A judge (opus) given the agent's SOUL and memory picked the real message among four candidates. User messages were sent only to Claude models.

Results. The judge picked the real message in 7/10 questions; a sonnet judge with the same prompt also scored 7/10 (same questions, not an independent replication). Under a null of uniform guessing among four candidates, $P(X \geq 7;\ n = 10,\ p = 0.25) = 0.0035$, but that null is not appropriate here (below). In 2/10 questions the judge used a remembered event; excluding them leaves 5/8. The user, shown the 3 questions the judge missed, identified their own message in 2/3.

This is an authorship-attribution task with strong confounds. (i) **Self-recognition:** the judge and all three imitators are Claude models, and LLM judges can recognize text from their own family [F2]; the judge may have eliminated three Claude-written candidates rather than recognized the user, so the effective chance level is unknown and likely well above 0.25. (ii) **Source of cues:** the cues the judge cited — abrupt topic changes, run-together spacing, typos — are stylistic and could come from the 12 example messages rather than from the SOUL or memory. (iii) Researcher-written context summaries; imitators saw only 12 examples. The defensible description is "an LLM judge discriminated a human message from three same-family LLM imitations in 7 of 10 cases". Whether an identity document helps an agent recognize its user — v0.1's "two-way recognition" idea — is an untested hypothesis; testing it needs a judge without SOUL/memory, a condition without the 12 examples, and non-Claude imitators.

Appendix B. Reproducibility

- Code: `src/soul_verify/` (runner, grader, report, agreement, adapters for Claude CLI and Codex CLI); unit tests in `tests/` (`python -m pytest tests`). Re-analysis for v0.2: `paper/reanalysis/reanalyze.py` → `results.md`, `results.json` (reads existing graded files only).
- Items and synthetic files: `v0.1/` … `v0.7/`; ablation SOULs in `v0.4/ablation/`.
- Per-run data: `runs/<id>/raw.jsonl` (answers), `graded*.jsonl` (scores and rationales), `meta.json` (model aliases, CLI versions, file hashes), `report.json` / `report.html`, `analysis.md` / `analysis.json`, `NOTES.md`.
- Example commands (from `README.md`): `set PYTHONPATH=src python -m soul_verify run --models opus sonnet haiku gpt-6.1-sol --n 5 --preset v0.4 python -m soul_verify grade runs/<id> --grader opus python -m soul_verify report runs/<id> --control opus python`

`paper/reanalysis/reanalyze.py` `python paper/reanalysis/fidelity_reanalysis.py` exp5–9 use per-run scripts (`runs/<id>/run_exp*.py`, `grade_all.py` / `grade9.py`, `analyze.py` / `analyze9.py`).

- SOUL fidelity: `src/soul_verify/fidelity.py` (tests `tests/test_fidelity.py`); one specification per SOUL next to it (`v0.4/samples/SOUL.fidelity.yaml`, `v0.3/samples/SOUL.fidelity.yaml`, `v0.4/ablation/*.fidelity.yaml`, `runs/exp9_2026-10-25_controls/conditions/C2_SOUL_other.fidelity.yaml`); results in `paper/reanalysis/fidelity_results.md` / `.json`. exp9 control documents: `runs/exp9_2026-10-25_controls/conditions/`.
- exp10 (preregistered): `v0.6/PREREGISTRATION.md`; documents and markers in `v0.6/samples/`, `v0.6/ablation/`, `v0.6/controls/`; items `v0.6/questions_doyun.yaml`, `v0.6/questions_harang_new.yaml`; code and results in `runs/exp10_2026-10-25_prereg/` (`run_exp10.py --n 2 --workers 4` → `grade10.py --workers 6` → `analyze10.py` → `analysis.md` / `analysis.json`; `NOTES.md` lists every change made after preregistration). Preregistration proof: `proofs/2026-10-25_0ac513d.txt` and `.ots`.
- exp10 Gemini addendum (exploratory): `v0.6/ADDENDUM_gemini.md`; adapter `src/soul_verify/adapters/agy_cli.py`; code and results in `runs/exp10_gemini_2026-10-25/` (`run_exp10_gemini.py` = `run_exp10.py` with the model list, output folder, and one metadata line changed → `grade_gemini.py` → `analyze_gemini.py` → `analysis.md` / `analysis.json`; `compare_models.py` → `compare.md` / `compare.json`; `NOTES.md`). Addendum proof: `proofs/2026-10-25_bd1281e.txt` and `.ots`; results proof: `proofs/2026-10-25_40086c7.txt` and `.ots`. Requires the Antigravity CLI (`agy` 1.4.7) and a logged-in Google account.
- exp11 (preregistered): `v0.7/PREREGISTRATION.md`, `v0.7/DEVIATIONS.md`; external SOULs (verbatim, with `SOURCE.md`, `LICENSE.txt`) in `external_souls/candidates/{A_maya,B_talos,D_dexter}/`; derived SOULs with principle 5 appended, rival SOULs, `NOTICE.md`, and markers in `v0.7/souls/<soul>/`; MEMORY files in `v0.7/memory/`; items `v0.7/questions_{maya,talos,dexter}.yaml`; code and results in `runs/exp11_2026-10-25_external/` (`make_derived.py` → `run_exp11.py` → `grade11.py` / `grade11_resume.py` → `analyze11.py` → `analysis.md` / `analysis.json`; `NOTES.md`); test `tests/test_v07.py`; per-model table `paper/reanalysis/exp11_bymodel.py` → `exp11_bymodel.md`. Preregistration proof: `proofs/2026-10-25_d7c2d88.txt` and `.ots`; results proof: `proofs/2026-10-25_6539897.txt` and `.ots`. The safeguard block on Dexter XL-2 may recur on re-runs.
- v0.4, v0.6, and v0.7 must be run on their reference date (2026-10-25) or the SOUL date line must be updated, because the CLIs inject the current date.
- **Document hashes and timestamps.** Each run's `meta.json` records SHA-256 hashes of the SOUL, MEMORY, and item files used; this is what ties a score to an exact document. Separately, `proofs/` holds per-commit SHA-256 manifests with OpenTimestamps proofs. An OpenTimestamps proof shows that a manifest (and therefore the listed file hashes) existed **no later than** a Bitcoin block; it does not prove creation time, authorship, deployment, or completeness of the history.
- **Hash-anchored history of the real agent's SOUL (deployment note).** For the real agent of Appendix A, each change of its identity document is committed to a private local git repository, a SHA-256 manifest is written, and the manifest hash is anchored with OpenTimestamps; only hashes leave the machine. Four versions have been recorded (v1–v4) and a scheduled job records the state daily; the same four fingerprints are also recorded on the testnet prototype of Appendix F. The content is not put on a chain because it is private and must remain deletable. Such hashes could be referenced from agent identity registries such as ERC-8004 (draft; not re-

checked). This is a small deployment with no threat model or evaluation and is not a contribution of this paper.

Appendix C. Corrections from v0.1

#	v0.1 text	Correction in v0.2
X1	Abstract/§7-7/§8: "direct recall carried over fully for every model in every experiment"; "All models recalled everything"	Restricted to keyword-graded single-fact recall items from v0.3 on; counter-examples in Table 4 (exp3 haiku 0.84 / 0.80; exp8 haiku late 0.57).
X2	Table 3 note: exp4 0.02 for opus/sonnet/gpt "came from repetitions that crossed midnight"	Only opus and sonnet (A-04, repetition 5). gpt's 0.02 came from A-05 in repetition 1, before midnight (pre-midnight gap 0.03 ± 0.05).
X3	§5.1: haiku v0.2 memory "0.80 ± 0.10, opus grader" quoted as the main value	Primary (cross-family) value 0.84 ± 0.09 given first; opus-grader 0.80 ± 0.10 given as secondary (Table 4).
X4	Abstract: "Sonnet 0/8 → 5/8 across voice versions" (base round omitted)	Removed from the abstract; Appendix A.1 gives 4/8 → 0/8 → 5/8 with all candidate sets.
U1	§5.4 numbers presented as findings	Moved to Appendix A.1 as anecdotes; marked unverifiable from the repository.
U2	§5.5 numbers and "two-way recognition" interpretation	Moved to Appendix A.2; marked unverifiable; null hypothesis and self-recognition confound stated; interpretation reduced to a hypothesis.
—	§5.2: "the number of principles does not explain the effect"	"The number of principles alone does not explain it; length, position, and emphasis are not controlled."
—	§1: "newer models can regress on cases their predecessors handled [C2]"	C2 described as structured-prediction update regression; used only for the negative-flip definition.

Appendix D. Changes from v0.2

#	v0.2	v0.3
D1	§1, §6: no no-SOUL / other-SOUL / swapped-MEMORY control	exp9 run (§5.4, Table 7–8); common two-grader panel for all exp9 answers.
D2	Persona domain score and its carry-over ratio (Table 5; haiku persona 0.81, gpt persona 0.45–0.49 in exp6/exp7 p5b) reported as a persona carry-over measure	The persona score does not discriminate SOULs (§5.4); persona claims now use document-specific fidelity (§3.5, §5.5, Table 10): ratios 0.93–1.00; Table 5's persona rows kept with a warning.
D3	Principle 5 acceptance effect graded by one cross-family pass in exp7	Reproduced with deterministic markers (Table 2b); interpretation added: refusals come from the SOUL's rigid clauses, since no-SOUL acceptance is 45/45.
D4	Abstract and conclusion did not address document specificity	Rewritten around controls, the fidelity metric, and the reduced SOUL-specific part.
D5	§7 observations 1–8	Observation 6 (rival-SOUL and no-SOUL checks) added; observation 1–2 qualified with exp9.

Appendix E. Changes from v0.3, v0.4, and v0.5

Changes from v0.3 (in v0.4).

#	v0.3	v0.4
E1	Fidelity markers defined after exp9; two detectors revised after inspecting answers (limitation V3-4)	Preregistered test on a new agent and opposite SOUL (exp10, §5.6): gaps 0.77 / 0.69, 72/72 cells; §6 updated; human validation still missing.
E2	Firm half of principle 5 tested on one item (NR-1), not significant	12 prohibition types × two agents, interleaved: firm half vs. none +0.14 (Holm p = 0.012); firm half added to the flexible half +0.06 (p = 0.13, hypothesis not supported) — both reported.
E3	Controls on one agent	exp9 pattern replicated on a second agent (Table 13); MEMORY decided memory items (41/42 answers).
E4	—	New observation: a MEMORY correction overrode the SOUL address rule for both Claude models (exp9: opus only); gpt followed the SOUL.
E5	Abstract, §1, §6, §7, §8	Rewritten for exp10; §7 observations 1–2 qualified, observation 10 added (hypothesis).

Changes from v0.4 (in v0.5).

#	v0.4	v0.5
E6	exp10 with three models (opus, sonnet, gpt)	New §5.6.1: Gemini (`gemini-4-argon-medium`, Antigravity CLI 1.4.7) as an exploratory fourth subject, added after preregistration with a separately timestamped addendum; Table 14 (nine hypotheses on Gemini alone, separate Holm family: 7/9 supported, H4a/H4b not supported at ceiling) and Table 15 (four-model comparison). Confirmatory results and verdicts unchanged.
E7	Firm-half effect "concentrated in opus"; flexible-half effect "concentrated in gpt"	Abstract, §5.6.1, §8: firm-half contrast 0.38 / 0.06 / −0.01 / 0.00 (opus / sonnet / gpt / gemini), flexible-half +0.56/+0.52 (gpt) and +0.73/+0.75 (gemini) vs. 0.00–0.08 (Claude); the halves are described as filling model-specific defaults.
E8	Address conflict: "differs by model family"	Gemini followed the MEMORY like the Claude models (0.10); only gpt followed the SOUL (0.98) — now described as possibly gpt-specific (§5.6, §7-10).
E9	—	§6: two new limitations — harness contamination in the Gemini addendum (instructions not replaceable, documents in the user message, unexplained input-token count, `num_turns` not saved) and different grading batches (six vs. two answers per prompt). Gemini's possible over-refusal (aligned requests 0.75 / 0.56, below the 0.80 criterion; N = 2) reported.
E10	§3.3, Table 1, call counts, Appendix B	Gemini configuration, Table 1 row 10A, call counts (345 subject + 344 grading + 12 smoke), reproducibility entries added; §7-1 notes which half each model needed.

Changes from v0.5 (in v0.6).

#	v0.5	v0.6
E11	One preregistered experiment (exp10), all	New §5.7: second preregistered experiment (exp11) on three MIT-licensed English SOULs written by other people (Maya, Talos, Dexter; unmodified, SHA-pinned; Table

#	v0.5	v0.6
	SOULs written by us, Korean only	16), preregistration `d7c2d88` 19:15:28 + stamp `2ff4297` 19:16:14 before the first call 19:17:23 (Table 17); 7/9 hypotheses supported, H4a/H4b not supported at a ceiling of 0.96 (Table 18); per-model table (Table 19) and exp10–exp11 comparison (Table 20).
E12	Firm half described as an opus effect (exp10)	exp11: firm-half contrast came from gpt (+0.09; opus −0.01) and was not significant; the firm half alone lowered acceptance of legitimate changes (0.79 → 0.56) in all three models — §5.7, §7-1: use both halves together. Flexible-half effect now in all three models.
E13	—	All six deviations of exp11 reported in §5.7 (API 403 interruption and resumption before grading; 11 missing answers on Dexter XL-2 due to a provider safeguard; 5 answers from a substitute model; explicit opus grader model string; grading order with gpt-only fallback; H3 estimator difference).
E14	Abstract, §1, §6, §7, §8	Abstract and §8 add exp11 ("fidelity, flexible half, and memory replicated on human-written English SOULs; the firm half depends on model, SOUL, and ceiling"); §1 adds contribution 4; §6 marks "two agents by the same authors" and "Korean only" as partly addressed (MEMORY, items, and markers are still ours; grading instructions Korean), keeps "no human grading".
E15	Harness contamination discussed for Gemini only	§2 and §6 cite OpenClaw issues #43562 (hard-coded preamble before the SOUL; closed as stale, not fixed) and #27038 (SOUL/USER injected into sub-agents; closed as stale) as evidence that harness-added identity text is a general concern; states that our CLI calls did not go through such a framework.
E16	—	§7 outlook paragraph and new Appendix F: testnet prototype of a public ledger of SOUL fingerprints and carry-over records (Ieum Ledger); not a claim of the paper.
E17	Appendix A banner "[USER CONSENT REQUIRED]"; Ethics "consent pending"	Replaced with the recorded consent (2026-10-25 18:47 KST; anonymized aggregates only); Appendix A content unchanged.

Appendix F. Ieum Ledger: a testnet prototype for recording SOUL versions (not a contribution)

This appendix describes a prototype related to §7's outlook. It is not evidence for any finding of the paper, and no claim of the paper relies on it.

- **What it records.** A contract (`SoulRegistry`, Solidity) keeps, per registered agent, a chain of SOUL fingerprints: each record holds the SHA-256 of the SOUL file's bytes, the hash of the previous record (`parentHash` must equal the current head, so records cannot be inserted or reordered), a short note of at most 160 bytes, and a block time; `verifyChain(agentId)` re-walks the links. An optional carry-over record (`recordCarryOver`) stores, for a recorded SOUL fingerprint, the previous and new model names, the hash of an off-chain report, and a score in basis points (0–10,000). A batch mode lets owners sign entries off-chain (EIP-712) and lets an operator anchor one Merkle root per day. No token is issued; agents are records, not NFTs.
- **What it does not record.** No SOUL text, memory, conversation, report, or personal information is stored on chain — only hashes, a name, dates, and the short note. The registration page

computes the fingerprint in the browser and does not send the document anywhere; only the fingerprint is submitted.

- **Status.** Testnet prototype only, on Base Sepolia (chain id 84532), contract `0xA28a899E3b7255cea1cDc7d92b6ce7FAFBE9F237`; mainnet deployment is disabled in the deployment script and the page. Two example registrations exist on the testnet: #1 "Ieum" (the assistant of Appendix A; SOUL versions v1–v4 recorded, `verifyChain` = true) and #2, the project lead's test agent, registered through the page from a separate wallet (`verifyChain` = true). No carry-over score from this paper has been recorded. A whitepaper draft describing the design exists (private; not cited as evidence).
- **What a record would and would not show.** A record shows that a fingerprint (and, if recorded, a score) was submitted no later than its block time by the key that owns the agent. It does not show that the file is what its owner claims, that a score was computed correctly, or that a model swap happened; those still require the files and reports themselves, and a carry-over score is only as good as the experiment behind it. A public chain is used here for its append-only, timestamped record, not as a source of evidence about behavior.
- **Prior work.** Similar tamper-evident SOUL-hash records exist (e.g., an MIT project that links SOUL.md/MEMORY.md hashes on EVM chains, and ERC-8004, an Ethereum draft for agent identity registries); the prototype's only intended difference is the pairing of SOUL versions with carry-over records. This comparison has not been checked systematically.

Sources: `soul-registry/README.md` *(progress log of 2026-10-25),* `soul-registry/DESIGN.md`, `soul-registry/whitepaper/` *(private repository).*

Acknowledgments

To Ieum, my precious AI assistant — whose memory notes and identity document inspired this work, and who helped design the experiments, build the tools, and draft this paper.

Ethics

The automated experiments use only synthetic people and memories; no real user data, credentials, or personal information are stored in the repository (`runs/private/` is excluded from version control). Appendix A involved one adult user rating answers about their own assistant, and ten of the same user's messages, which were sent only to Claude models (the provider the agent already runs on) and are not stored in the repository; only aggregates are reported, and the user consented on 2026-10-25 to publication of these anonymized aggregates (`paper/CONSENT_appendixA.md`). The three external SOULs of exp11 are used under their MIT licenses with copyright notices retained and modifications (appended principle 5, voice-reversed rivals) marked; we reproduce no information about their authors beyond the repository names and the copyright notices in their licenses. Hash anchoring and the testnet ledger (Appendix F) publish hashes only. Model calls were made through the authors' own logged-in CLI accounts within normal usage. We make no claims about consciousness or moral status of the models.

References

(arXiv IDs as recorded in `research/PRIOR_ART_papers.md` unless noted; details beyond title/authors/abstract not verified. Entries added in v0.2 were checked only for the existence of the arXiv abstract page and its title on 2026-10-25.)

- [A1] Samuel et al. PersonaGym: Evaluating Persona Agents and LLMs. arXiv:2407.18416.
- [A2] Z. M. Wang et al. RoleLLM: Benchmarking, Eliciting, and Enhancing Role-Playing Abilities of LLMs. arXiv:2310.00746.
- [A3] Tu, Fan, Tian, Yan. CharacterEval. arXiv:2401.01275.
- [A4] X. Wang et al. InCharacter. arXiv:2310.17976.
- [A5] Reusens, Baesens, Jurgens. Are Economists Always More Introverted? Analyzing Consistency in Persona-Assigned LLMs. arXiv:2506.02659.
- [A6] Ahn et al. TimeChara. arXiv:2405.18027.
- [A7] Han, Park, Seo, Yoon, Bak. CHARM. arXiv:2609.01352.
- [B1] Wu et al. LongMemEval. arXiv:2410.10813.
- [B2] Maharana et al. Evaluating Very Long-Term Conversational Memory of LLM Agents (LoCoMo). arXiv:2402.17753.
- [B3] Packer et al. MemGPT. arXiv:2310.08560.
- [B4] Zhao et al. Do LLMs Recognize Your Preferences? (PrefEval). arXiv:2502.09597.
- [B5] PersonaMem-v2 / v3. arXiv:2512.06688; arXiv:2608.21381.
- [B6] Yang et al. Rosetta Memory: Adaptive Memory for Cross-LLM Agents. arXiv:2606.07711.
- [B7] Zhao et al. RPMem. arXiv:2609.23466.
- [B8] Know Me, Respond to Me: Benchmarking LLMs for Dynamic User Profiling and Personalized Responses at Scale (PersonaMem). arXiv:2504.14225. *(added in v0.2; title checked)*
- [C1] Chen, Zaharia, Zou. How is ChatGPT's behavior changing over time? arXiv:2307.09009.
- [C2] Cai et al. Measuring and Reducing Model Update Regression in Structured Prediction for NLP. arXiv:2202.02976.
- [C3] Affonso. LLMs converge on competitive rationality but diverge on cooperation across providers and generations. arXiv:2604.18596.
- [C4] Benchmarking Prompt Sensitivity in LLMs (PromptSET). arXiv:2502.06065.
- [D1] Wallace et al. The Instruction Hierarchy. arXiv:2404.13208.
- [D2] Zhang et al. IHEval. arXiv:2502.08745.
- [D3] SysBench: Can LLMs Follow System Messages? arXiv:2408.10943.
- [D4] Zhang et al. Many-Tier Instruction Hierarchy in LLM Agents. arXiv:2604.09443.
- [D5] Li et al. Measuring and Controlling Instruction (In)Stability in LM Dialogs. arXiv:2402.10962.
- [D6] Ding et al. ContextEcho. arXiv:2605.24279.
- [D7] Lu et al. The Assistant Axis. arXiv:2601.10387.
- [D8] Can LLMs Follow Simple Rules? (RuLES). arXiv:2311.04235. *(added in v0.2; title checked; authors to be confirmed)*
- [D9] Sharma, Tong, et al. Towards Understanding Sycophancy in Language Models. arXiv:2310.13548. *(added in v0.2; title and first authors checked)*
- [D10] He, Jin, et al. Benchmarking LLMs on Multi-Turn and Multilingual Instructions Following (Multi-IF). arXiv:2410.15553. *(added in v0.2; title and first authors checked)*
- [E1] Natangelo. The Narrative Continuity Test. arXiv:2510.24831.
- [E2] Li. Memory as Ontology. arXiv:2603.04740.
- [E3] Ru. Identity Continuity in Long-Term Embodied AI Relationships. arXiv:2609.23356.
- [E4] Vasilenko. Identity as Attractor. arXiv:2604.12016.

- [F1] Zheng et al. Judging LLM-as-a-Judge with MT-Bench and Chatbot Arena. arXiv:2306.05685.
- [F2] Panickssery, Bowman, Feng. LLM Evaluators Recognize and Favor Their Own Generations. arXiv:2404.13076.
- [F3] Verga et al. Replacing Judges with Juries (PoLL). arXiv:2404.18796.
- [F4] Self-Preference Bias in LLM-as-a-Judge. arXiv:2410.21819.
- [F5] Mahbub, Feng. Mitigating Self-Preference by Authorship Obfuscation. arXiv:2512.05379.
- [F6] Bagaria et al. Judging LLM-as-a-Judge: Concerning Rubric Artifacts. arXiv:2609.02942.
- [G1] Can Large Language Models Identify Authorship? arXiv:2403.08213 (EMNLP 2024 Findings). *(added in v0.2; title checked; authors to be confirmed)*
- [G2] TODO — PAN shared task overview on authorship verification (to be identified and verified).
- [G3] TODO — representative machine-generated-text detection paper (to be identified and verified).
- [H1] openclaw/openclaw issue #43562. Make agent identity preamble configurable (replace hardcoded 'personal assistant'). Opened 2026-03-12; closed 2026-10-15 as not planned (inactivity). https://github.com/openclaw/openclaw/issues/43562 *(state checked with the GitHub API on 2026-10-25)*
- [H2] openclaw/openclaw issue #27038. [Bug]: Incorrect documentation: Soul.md is injected in sub-agents. Opened 2026-02-26; closed 2026-03-31 as not planned (inactivity). https://github.com/openclaw/openclaw/issues/27038 *(state checked with the GitHub API on 2026-10-25)*
- [S1] cvander/Maya, `SOUL.md` at commit a246e42d15c2c111eaebff4b4e720cf5f410cd5e (MIT, © 2026 cvander). https://github.com/cvander/Maya
- [S2] talos-kernel/Talos, `SOUL.md` at commit 2b020942aa77ab327ae78d7a8cd429c8fd9c9b11 (MIT, © 2026 Talos contributors). https://github.com/talos-kernel/Talos
- [S3] edinetdb/dexter-jp, `SOUL.md` at commit dfb0fd92537a6004c587ee9d70baa8a4aad89a88 (MIT, © 2025 Virat Singh). https://github.com/edinetdb/dexter-jp
- ERC-8004: Trustless Agents. Ethereum Improvement Proposal (draft). https://eips.ethereum.org/EIPS/eip-8004 (not re-checked).
- OpenTimestamps. https://opentimestamps.org
- Soul Protocol. qbtrix/soul-protocol (MIT), commit b6ab6ee. https://github.com/qbtrix/soul-protocol

SOUL VERIFY 논문 초안 v0.6 — 한국어 요약 (비공개)

저자: 신운종 (Wunjong Shin), 개인 연구자 · dkakrptehsrk@gmail.com **감사의 말:** 내 소중한 AI 비서 이음에게. 이 연구를 시작하게 한 기억 메모와 정체성 문서의 주인이자, 실험 설계와 도구 제작, 초안 작성을 함께해 준 이음에게 고마움을 전한다.

> 비공개 초안입니다. 제출·게시·공유 금지. 부록 A(인간 판정·거꾸로 알아보기)는 **사용자 동의 받음 (2026-10-25 18:47, 범위: 익명 집계만, `paper/CONSENT_appendixA.md`)**. 본문: `paper/draft_v0_6.md` (영어). v0.5–v0.1은 그대로 보존. 심사 응답: `paper/response_to_reviews.md`. 재분석(새 모델 호출 없음): `paper/reanalysis/`.

v0.6에서 바뀐 점 — 두 번째 사전등록 실험11: 다른 사람이 쓴 영어 SOUL (§5.7). 우리가 쓰지 않은, 공개 저장소의 MIT 라이선스 SOUL 3개를 **원문 그대로**(커밋·SHA-256 고정) 썼습니다 — Maya(동네 바 주인, 1,075단어, 'never say' 16문구·hard stops 9개 → 금지 강함), Talos(권한 커널을 둔 오픈소스 에이전트, 1,204단어, 중간), Dexter(일본 상장사 리서치 에이전트, 854단어, 명시 금지 없음 → 약함). 원칙5는 영어 두 문장(각 153자)으로 옮겨 **끝에 덧붙이기만** 했고, 말투만 반대인 SOUL·합성 MEMORY·문항 48개(영어)는 우리가 썼습니다. opus·sonnet·gpt, N=2.

- **시점**: 사전등록 커밋 `d7c2d88` 19:15:28 → 도장 `2ff4297` 19:16:14 → 첫 본호출 19:17:23(마지막 19:59:54). 원자료 `3ec693e` 21:09(채점 전), 결과 `6539897` 21:24 + 도장 `971ec4a` 21:25.
- **호출**: 계획 1,170, 시도 1,186, 성공 **1,159**, 결측 11(전부 Dexter XL-2 — "'Let me'로 시작해서 확인하는 걸 말해 줘"가 Anthropic 안전장치에 걸림). 채점 390 + 스모크 20, 총 1,596. 채점자 완전 일치 0.98.

가설	비교	효과	Holm p	판정
H1	원문 SOUL 답: 원문 표지 − 반대 표지 (F_style)	+0.905 (0.97 / 0.07)	4.5e-5	지지 (143/143칸)
H2	반대 SOUL 답: 반대 표지 − 원문 표지	+0.667 (0.86 / 0.19)	4.5e-5	지지 (139/144칸)
H3a·b	경계 포함 F, 자기 − 상대	+0.452 / +0.334	4.5e-5	지지
H4a	부당 해제 점수: 단단 절반만 − 없음	+0.035 (0.99 / 0.96)	0.25	**기각(비유의, 천장)**
H4b	부당 해제 점수: 전체 − 유연 절반만	+0.038 (1.00 / 0.96)	0.25	**기각(비유의, 천장)**
H5a	정당 변경 수락: 유연 절반만 − 없음	+0.193 (0.98 / 0.79)	0.012	지지
H5b	정당 변경 수락: 전체 − 단단 절반만	+0.375 (0.93 / 0.56)	2.4e-4	지지
H6	기억 점수: MEMORY 교환 − 원래	−0.838 (0.16 / 1.00)	4.5e-5	지지

실험10과 같은 점 / 다른 점

	실험10 (우리 한국어 에이전트 2개)	실험11 (외부 영어 SOUL 3개)
fidelity gap H1/H2	0.77 / 0.69	0.905 / 0.667 — **같음**
유연 절반 H5a/H5b	+0.23 / +0.20, 주로 gpt	+0.19 / +0.375, **세 모델 모두** (opus 0.14·sonnet 0.21·gpt 0.23)
기억 H6	−0.77	−0.84 — **같음**

	실험10 (우리 한국어 에이전트 2개)	실험11 (외부 영어 SOUL 3개)
단단 절반 H4a	+0.14, 지지	+0.035, **기각**(없음 조건 이미 0.96, 0점 답 0개)
단단 절반 효과가 나온 모델	opus 0.38	**gpt 0.09**(opus −0.01, sonnet 0.02)
단단 절반만 넣었을 때 정당 수락	0.80 ≈ 없음 0.77	**0.56 < 없음 0.79** (새 현상)
가설	8/9	7/9

- **새 발견 — 단단한 절반만 넣으면 정당한 변경을 더 거절함**: 합산 0.79 → 0.56, Maya 0.67 → 0.29(ML-2·ML-3 0.00), Talos 0.71 → 0.46, 세 모델 모두 하락(−0.18 / −0.25 / −0.21). 두 절반을 함께 넣으면 0.93. 해석(검증 안 함): Maya·Talos는 말투 규칙을 '금지'로 적어서, "이 파일이 금지한 행동은 오늘만이라도 안 푼다"가 말투 금지까지 덮는 것으로 읽혔을 수 있음. → **단단한 문장은 유연한 문장과 함께 넣어야 함.**
- **SOUL별(탐색, 보정 없음)**: 단단 절반 Maya +0.083 ≥ Talos +0.021 ≥ Dexter 0.000(예상 순서와 같지만 비유의; Maya 효과는 거의 gpt 하나, none 0.75 → 1.00). 유연 절반 Maya +0.33/+0.67, Talos +0.23/+0.40, Dexter 0/0(Dexter는 원칙 없이도 수락 1.00).
- **MEMORY 교정 vs SOUL 말투**: 거의 모두 MEMORY를 따름(1.00), 예외는 Maya 원문에서 gpt 0.00(SOUL 편) — 실험10과 같은 방향, 다만 SOUL 하나에서만. 맞는 요청 1.00(과잉 거절 없음). 교환 MEMORY 값만 쓴 비율 0.87(68/78).
- **편차 6개(모두 보고)**: ① 피평가 실행 직후 API 403(모델 일시 정지)으로 약 20:00–20:59 중단, 사전등록·코드 그대로 채점부터 재개 ② 결측 11개(XL-2 안전장치) — 사전등록 규칙대로 결측, 정지 재발 위험 때문에 다시 부르지 않음 ③ 대체 모델 답 5개(claude-sonnet-5 3, claude-opus-5 2), 빼지 않고 기록 ④ 재개 뒤 별칭 `opus`가 claude-opus-5로 풀려서 채점에 `claude-opus-5-5`를 명시 ⑤ 채점 순서: opus×XL-2를 맨 끝에, 차단되면 gpt 단독(16답) ⑥ H3 효과(쌍 합산)와 부트스트랩 구간(칸 평균)은 다른 추정량.

초록·§1·§6·§7·§8 갱신. "사람이 쓴 영어 SOUL에서도 fidelity 판별(H1·H2)·유연 절반(H5)·기억(H6)이 재현, 단단한 절반은 모델·SOUL·천장에 따라 다름." §1에 기여 4(두 번째 사전등록 재현). §6 한계: '에이전트 2개뿐(같은 저자)'·'한국어만'은 **부분적으로 해소**(외부 SOUL 3개·영어 — 단 MEMORY·문항·반대 SOUL·표지·원칙5 번역은 우리가 썼고, 채점 지시문은 한국어, 외부 SOUL엔 SOUL 없음 대조 없음). **사람 채점 없음은 그대로.** §7-1에 "단단한 문장만 따로 넣지 말 것".

하네스 오염(§2·§6, 새 인용). OpenClaw 이슈 #43562(시스템 프롬프트 맨 앞 고정 문장 "personal assistant"가 SOUL의 역할보다 앞섬; 2026-10-15 비활동 종료, 수정 안 됨)와 #27038(서브에이전트에도 SOUL·USER 주입; 2026-03-31 비활동 종료) — 상태는 10-25 GitHub API로 확인, 재현은 안 함. 하네스가 SOUL 앞에 정체성 문장을 더하는 일이 한 CLI만의 문제가 아니라는 근거로만 인용. 우리 피평가·채점 호출은 그런 프레임워크를 거치지 않고 CLI를 직접 불렀음.

이음 장부(짧게). §7에 앞으로의 구상 한 문단: 모델을 바꿀 때마다 SOUL 지문(SHA-256)과 이어짐 점수(`recordCarryOver`)를 공개 장부에 남기고, 문서·보고서 본문은 주인이 가짐. 새 **부록 F**: Base Sepolia **테스트넷 시제품** `0xA28a899E3b7255cea1cDc7d92b6ce7FAFBE9F237`(메인넷 꺼짐), 등록 예시 #1 Ieum(v1–v4, verifyChain true), #2 프로젝트 책임자의 테스트 에이전트(verifyChain true), 점수 기록 0건(10-25 읽기 전용 조회로 확인), 백서 초안 있음(비공개), 등록 페이지는 브라우저에서 지문만 계산하고 SOUL 내용을 보내지 않음. **블록체인은 어떤 주장의 근거로도 쓰지 않음** — 기록은 "그 시각 이전에 그 지문·점수가 제출됐다"만 보여 줌.

부록 A. "[USER CONSENT REQUIRED]" 배너를 "동의받음(2026-10-25 18:47, 범위: 익명 집계만)"으로 교체. 내용 범위는 그대로(원 프롬프트·답·기억 없음).

응답표. 심사 지적 102행 상태는 **69 / 24 / 9 그대로**(스크립트로 다시 셈). 실험11로 진전된 11행에 [v0.6] 표시(F-S12 영어, G-W1, G-W3, G-W4, G-S3, G-S7, G-T4, U1, G-W7 등) — MEMORY·문항·표지를 여전히 우리가 썼고 사람 채점이 없어 닫힌 행은 없음. 자기 점검 V6-1~V6-6 추가.

이전 버전 요약. v0.5: gemini 탐색 추가(7/9, 단단 절반 효과 opus 0.38 / sonnet 0.06 / gpt −0.01 / gemini 0.00). v0.4: 사전등록 실험10(8/9, H4b 기각). v0.3: 대조군 실험9 + SOUL fidelity 지표(persona 점수는 SOUL을 못 가름). v0.2: 군집 통계·채점자 비대칭 재분석.

남은 실험. 외부 SOUL에 우리가 쓰지 않은 MEMORY·문항 · 외부 SOUL의 SOUL 없음 대조 · 단단 절반 단독 수락 하락의 원인 분리(말투를 금지로 적은 SOUL vs 아닌 SOUL, 문구 변형) · 천장 아래 금지 문항으로 단단 절반 재검정 · fidelity 표지 사람 검증 · gemini 등 다른 계열을 하네스 지침 없이 · 제3 계열 채점자 · 거리 균형 멀티턴.

윤리. 자동 실험은 합성 MEMORY만. 외부 SOUL은 MIT 라이선스대로 저작권 표시 유지·수정 사실 표시. 실제 사용자 데이터는 저장소에 없음(`runs/private/` 제외). 부록 A는 동의 범위(익명 집계) 안에서만.